


Chapter 3: Processes



Chapter 3: Processes

- ▶ Process Concept
- ▶ Process Scheduling
- ▶ Operations on Processes
- ▶ Interprocess Communication
- ▶ Examples of IPC Systems
- ▶ Communication in Client-Server Systems



Objectives

- ▶ To introduce the notion of a process -- a program in execution, which forms the basis of all computation
- ▶ To describe the various features of processes, including scheduling, creation and termination, and communication
- ▶ To describe communication in client-server systems



Process Concept

- ▶ An operating system executes a variety of programs:
 - ◆ Batch system – jobs
 - ◆ Time-sharing systems – user programs or tasks
- ▶ Textbook uses the terms *job* and *process* almost interchangeably
- ▶ Process – a program in execution; process execution must progress in sequential fashion
- ▶ A process includes:
 - ◆ program counter
 - ◆ stack
 - ◆ data section






The Process

- ▶ Multiple parts
 - ✦ The program code, also called **text section**
 - ✦ Current activity including **program counter**, processor registers
 - ✦ **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - ✦ **Data section** containing global variables
 - ✦ **Heap** containing memory dynamically allocated during run time
- ▶ Program is passive entity, process is active
 - ✦ Program becomes process when executable file loaded into memory
- ▶ Execution of program started via GUI mouse clicks, command line entry of its name, etc
- ▶ One program can be several processes
 - ✦ Consider multiple users executing the same program



Process in Memory




Process State

- ▶ As a process executes, it changes *state*
 - ✦ **new**: The process is being created
 - ✦ **running**: Instructions are being executed
 - ✦ **waiting**: The process is waiting for some event to occur
 - ✦ **ready**: The process is waiting to be assigned to a processor
 - ✦ **terminated**: The process has finished execution



Diagram of Process State





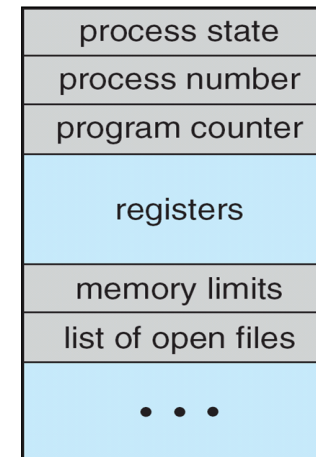
Process Control Block (PCB)

Information associated with each process

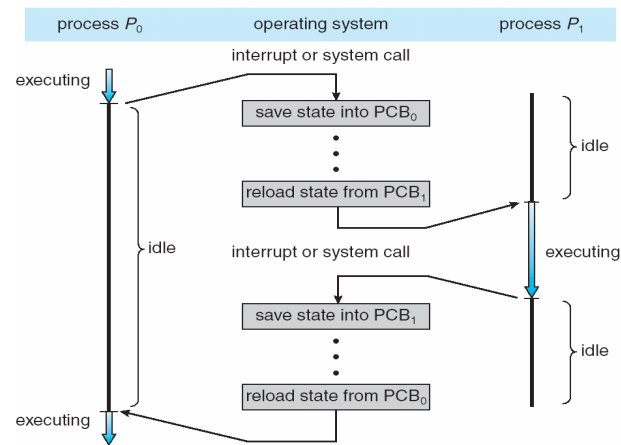
- ▶ Process state
- ▶ Program counter
- ▶ CPU registers
- ▶ CPU scheduling information
- ▶ Memory-management information
- ▶ Accounting information
- ▶ I/O status information



Process Control Block (PCB)



CPU Switch From Process to Process



Process Scheduling

- ▶ Maximize CPU use, quickly switch processes onto CPU for time sharing
- ▶ **Process scheduler** selects among available processes for next execution on CPU
- ▶ Maintains **scheduling queues** of processes
 - ✦ **Job queue** – set of all processes in the system
 - ✦ **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - ✦ **Device queues** – set of processes waiting for an I/O device
 - ✦ Processes migrate among the various queues






Process Representation in Linux

► Represented by the C structure `task_struct`


```

pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this pro */
    
```

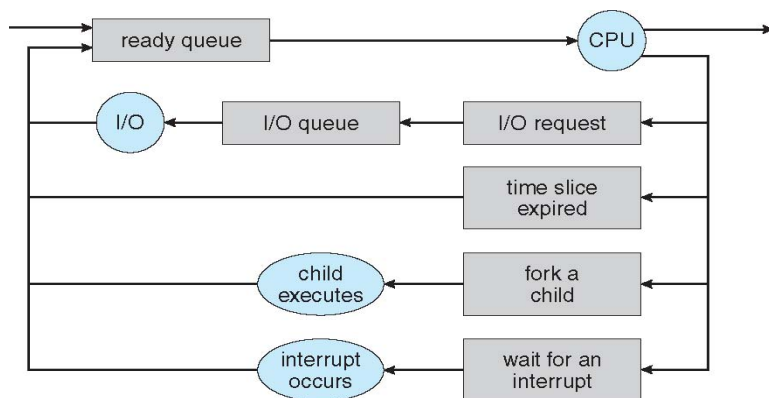
struct
struct



Ready Queue And Various I/O Device Queues



Representation of Process Scheduling



Schedulers

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU
 - ◆ Sometimes the only scheduler in a system



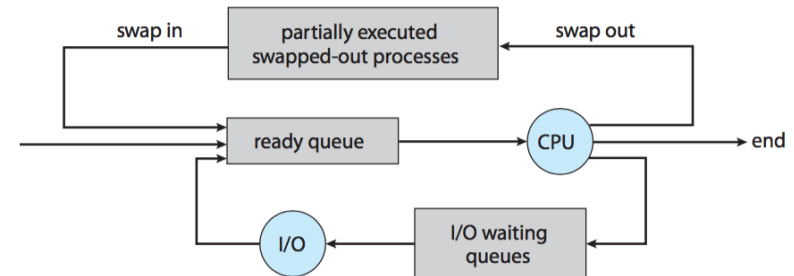


Schedulers (Cont.)

- ▶ Short-term scheduler is invoked very frequently (milliseconds) ⇒ (must be fast)
- ▶ Long-term scheduler is invoked very infrequently (seconds, minutes) ⇒ (may be slow)
- ▶ The long-term scheduler controls the *degree of multiprogramming*
- ▶ Processes can be described as either:
 - ♦ **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - ♦ **CPU-bound process** – spends more time doing computations; few very long CPU bursts



Addition of Medium Term Scheduling



Context Switch

- ▶ When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a **context switch**.
- ▶ **Context** of a process represented in the PCB
- ▶ Context-switch time is overhead; the system does no useful work while switching
 - ♦ The more complex the OS and the PCB -> longer the context switch
- ▶ Time dependent on hardware support
 - ♦ Some hardware provides multiple sets of registers per CPU -> multiple contexts loaded at once



Process Creation

- ▶ **Parent** process create **children** processes, which, in turn create other processes, forming a tree of processes
- ▶ Generally, process identified and managed via a **process identifier (pid)**
- ▶ Resource sharing
 - ♦ Parent and children share all resources
 - ♦ Children share subset of parent's resources
 - ♦ Parent and child share no resources
- ▶ Execution
 - ♦ Parent and children execute concurrently
 - ♦ Parent waits until children terminate






Process Creation (Cont.)

- ▶ Address space
 - ◆ Child duplicate of parent
 - ◆ Child has a program loaded into it
- ▶ UNIX examples
 - ◆ **fork** system call creates new process
 - ◆ **exec** system call used after a **fork** to replace the process' memory space with a new program



Process Creation



C Program Forking Separate Process

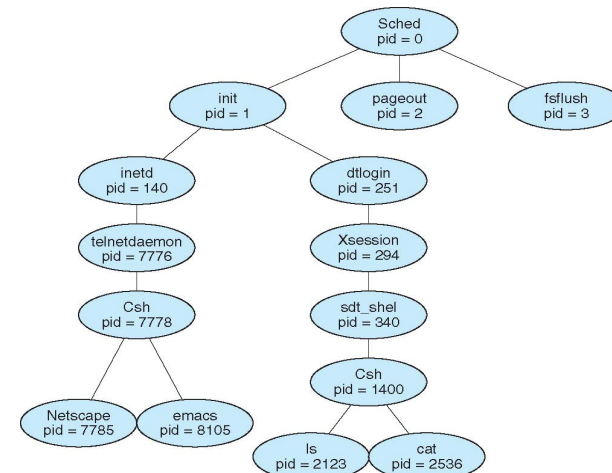
```

#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}

```



A Tree of Processes on Solaris





Process Termination

- ▶ Process executes last statement and asks the operating system to delete it (**exit**)
 - ✦ Output data from child to parent (via **wait**)
 - ✦ Process' resources are deallocated by operating system
- ▶ Parent may terminate execution of children processes (**abort**)
 - ✦ Child has exceeded allocated resources
 - ✦ Task assigned to child is no longer required
 - ✦ If parent is exiting
 - ▶ Some operating systems do not allow child to continue if its parent terminates
 - All children terminated - **cascading termination**

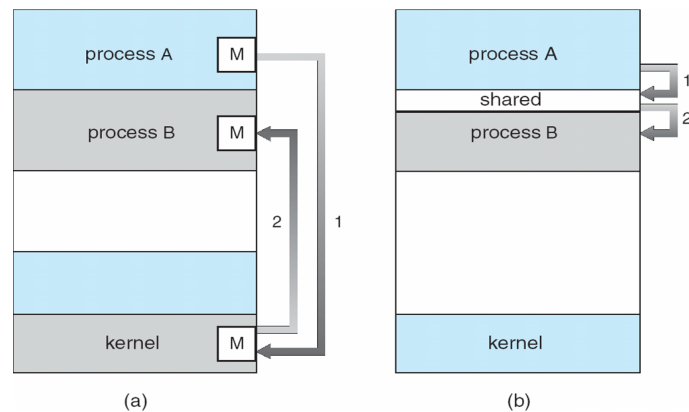


Interprocess Communication

- ▶ Processes within a system may be **independent** or **cooperating**
- ▶ Independent process cannot affect or be affected by the execution of another process
- ▶ Cooperating process can affect or be affected by other processes, including sharing data
- ▶ Reasons for cooperating processes:
 - ✦ Information sharing
 - ✦ Computation speedup
 - ✦ Modularity
 - ✦ Convenience
- ▶ Cooperating processes need **interprocess communication (IPC)**
- ▶ Two models of IPC
 - ✦ Shared memory
 - ✦ Message passing



Communications Models



Producer-Consumer Problem

- ▶ Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - ✦ *unbounded-buffer* places no practical limit on the size of the buffer
 - ✦ *bounded-buffer* assumes that there is a fixed buffer size





Bounded-Buffer – Shared-Memory Solution

► Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

► Solution is correct, but can only use BUFFER_SIZE-1 elements



Bounded-Buffer – Producer

```
item next_produced;
while (true) {
    /* Produce an item in next_produced*/
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing -- no free buffers */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```



Consumer

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing -- nothing to consume */
    next_consumed = buffer[out]
    ; /* consume an item from the buffer */
    out = (out + 1) % BUFFER_SIZE;
}
```



Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - ◆ **send**(message) – message size fixed or variable
 - ◆ **receive**(message)
- If *P* and *Q* wish to communicate, they need to:
 - ◆ establish a *communication link* between them
 - ◆ exchange messages via send/receive
- Implementation of communication link
 - ◆ physical (e.g., shared memory, hardware bus)
 - ◆ logical (e.g., logical properties)





Implementation Questions

- ▶ How are links established?
- ▶ Can a link be associated with more than two processes?
- ▶ How many links can there be between every pair of communicating processes?
- ▶ What is the capacity of a link?
- ▶ Is the size of a message that the link can accommodate fixed or variable?
- ▶ Is a link unidirectional or bi-directional?



Direct Communication

- ▶ Processes must name each other explicitly:
 - ♦ *send(P, message)* – send a message to process P
 - ♦ *receive(Q, message)* – receive a message from process Q
- Properties of communication link
 - ♦ Links are established automatically
 - ♦ A link is associated with exactly one pair of communicating processes
 - ♦ Between each pair there exists exactly one link
 - ♦ The link may be unidirectional, but is usually bi-directional



Indirect Communication

- ▶ Messages are directed and received from mailboxes (also referred to as ports)
 - ♦ Each mailbox has a unique id
 - ♦ Processes can communicate only if they share a mailbox
- ▶ Properties of communication link
 - ♦ Link established only if processes share a common mailbox
 - ♦ A link may be associated with many processes
 - ♦ Each pair of processes may share several communication links
 - ♦ Link may be unidirectional or bi-directional



Indirect Communication

- ▶ Operations
 - ♦ create a new mailbox
 - ♦ send and receive messages through mailbox
 - ♦ destroy a mailbox
- ▶ Primitives are defined as:
 - send(A, message)** – send a message to mailbox A
 - receive(A, message)** – receive a message from mailbox A





Indirect Communication

- ▶ Mailbox sharing
 - ✦ P_1 , P_2 , and P_3 share mailbox A
 - ✦ P_1 sends; P_2 and P_3 receive
 - ✦ Who gets the message?
- ▶ Solutions
 - ✦ Allow a link to be associated with at most two processes
 - ✦ Allow only one process at a time to execute a receive operation
 - ✦ Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.



Synchronization

- ▶ Message passing may be either blocking or non-blocking
- ▶ **Blocking** is considered **synchronous**
 - ✦ **Blocking send** has the sender block until the message is received
 - ✦ **Blocking receive** has the receiver block until a message is available
- ▶ **Non-blocking** is considered **asynchronous**
 - ✦ **Non-blocking send** has the sender send the message and continue
 - ✦ **Non-blocking receive** has the receiver receive a valid message or null



Buffering

- ▶ Queue of messages attached to the link; implemented in one of three ways
 1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits



Examples of IPC Systems - POSIX

- ▶ POSIX Shared Memory
 - ✦ Process first creates shared memory segment

```
segment_id = shmget(IPC_PRIVATE, size, S_IRUSR | S_IWUSR);
```
 - ✦ Process wanting access to that shared memory must attach to it

```
shared_memory = (char *) shmat(id, NULL, 0);
```
 - ✦ Now the process could write to the shared memory

```
sprintf(shared_memory, "Writing to shared memory");
```
 - ✦ When done a process can detach the shared memory from its address space

```
shmdt(shared_memory);
```





Examples of IPC Systems - Mach

- ▶ Mach communication is message based
 - ◆ Even system calls are messages
 - ◆ Each task gets two mailboxes at creation- Kernel and Notify
 - ◆ Only three system calls needed for message transfer
`msg_send()`, `msg_receive()`, `msg_rpc()`
 - ◆ Mailboxes needed for communication, created via
`port_allocate()`

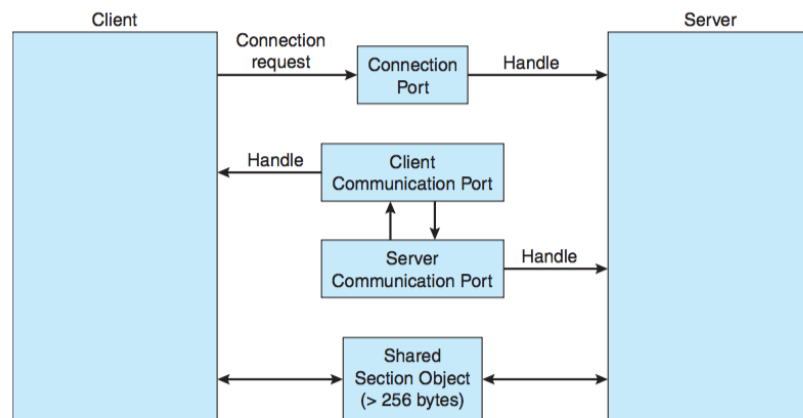


Examples of IPC Systems – Windows XP

- ▶ Message-passing centric via **Advanced local procedure call (ALPC)** facility
 - ◆ Only works between processes on the same system
 - ◆ Uses ports (like mailboxes) to establish and maintain communication channels
 - ◆ Communication works as follows:
 - ▶ The client opens a handle to the subsystem's **connection port** object.
 - ▶ The client sends a connection request.
 - ▶ The server creates two private **communication ports** and returns the handle to one of them to the client.
 - ▶ The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.



Local Procedure Calls in Windows XP



Communications in Client-Server Systems

- ▶ Sockets
- ▶ Remote Procedure Calls
- ▶ Pipes






Sockets

- ▶ A **socket** is defined as an *endpoint for communication*
- ▶ Concatenation of IP address and port
- ▶ The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- ▶ Communication consists between a pair of sockets



Socket Communication

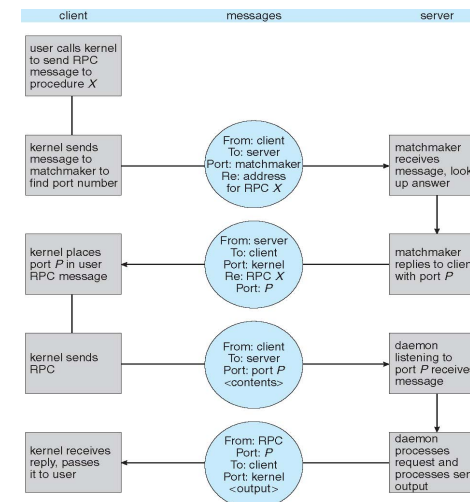


Remote Procedure Calls

- ▶ Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
- ▶ **Stubs** – client-side proxy for the actual procedure on the server
- ▶ The client-side stub locates the server and *marshalls* the parameters
- ▶ The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server



Execution of RPC






Pipes

- ▶ Acts as a conduit allowing two processes to communicate
- ▶ **Issues**
 - ✦ Is communication unidirectional or bidirectional?
 - ✦ In the case of two-way communication, is it half or full-duplex?
 - ✦ Must there exist a relationship (i.e. parent-child) between the communicating processes?
 - ✦ Can the pipes be used over a network?



Pipes



Ordinary Pipes

- ▶ **Ordinary Pipes** allow communication in standard producer-consumer style
- ▶ Producer writes to one end (the *write-end* of the pipe)
- ▶ Consumer reads from the other end (the *read-end* of the pipe)
- ▶ Ordinary pipes are therefore unidirectional
- ▶ Require parent-child relationship between communicating processes



Named Pipes

- ▶ Named Pipes are more powerful than ordinary pipes
- ▶ Communication is bidirectional
- ▶ No parent-child relationship is necessary between the communicating processes
- ▶ Several processes can use the named pipe for communication
- ▶ Provided on both UNIX and Windows systems



End of Chapter 3

