


Chapter 4: Threads



Chapter 4: Threads

- ▶ Overview
- ▶ Multithreading Models
- ▶ Thread Libraries
- ▶ Threading Issues
- ▶ Operating System Examples
- ▶ Windows XP Threads
- ▶ Linux Threads



Objectives

- ▶ To introduce the notion of a thread — a fundamental unit of CPU utilization that forms the basis of multithreaded computer systems
- ▶ To discuss the APIs for the Pthreads, Win32, and Java thread libraries
- ▶ To examine issues related to multithreaded programming



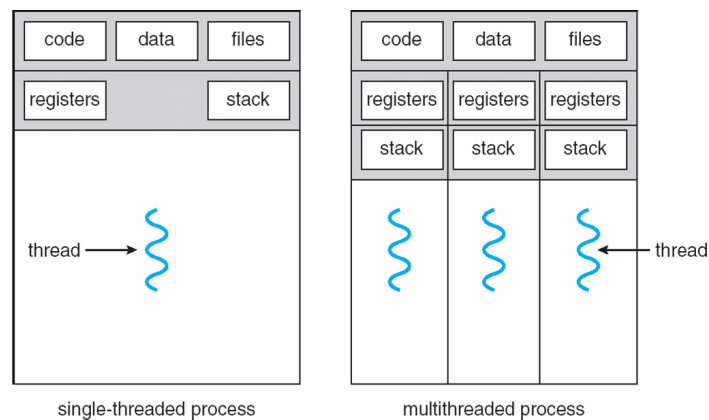
Motivation

- ▶ Threads run within application
- ▶ Multiple tasks with the application can be implemented by separate threads
 - ◆ Update display
 - ◆ Fetch data
 - ◆ Spell checking
 - ◆ Answer a network request
- ▶ Process creation is heavy-weight while thread creation is light-weight
- ▶ Can simplify code, increase efficiency
- ▶ Kernels are generally multithreaded





Single and Multithreaded Processes



Benefits

- ▶ Responsiveness
- ▶ Resource Sharing
- ▶ Economy
- ▶ Scalability




Multicore Programming

- ▶ Multicore systems putting pressure on programmers, challenges include:
 - ◆ Dividing activities
 - ◆ Balance
 - ◆ Data splitting
 - ◆ Data dependency
 - ◆ Testing and debugging




Multithreaded Server Architecture






Concurrent Execution on a Single-core System



Parallel Execution on a Multicore System



User Threads

- ▶ Thread management done by user-level threads library
- ▶ Three primary thread libraries:
 - ♦ POSIX **Pthreads**
 - ♦ Win32 threads
 - ♦ Java threads



Kernel Threads

- ▶ Supported by the Kernel
- ▶ Examples
 - ♦ Windows XP/2000
 - ♦ Solaris
 - ♦ Linux
 - ♦ Tru64 UNIX
 - ♦ Mac OS X





Multithreading Models

- ▶ Many-to-One
- ▶ One-to-One
- ▶ Many-to-Many




Many-to-One

- ▶ Many user-level threads mapped to single kernel thread
- ▶ Examples:
 - ♦ Solaris Green Threads
 - ♦ GNU Portable Threads



Many-to-One Model




One-to-One

- ▶ Each user-level thread maps to kernel thread
- ▶ Examples
 - ♦ Windows NT/XP/2000
 - ♦ Linux
 - ♦ Solaris 9 and later





One-to-one Model




Many-to-Many Model

- ▶ Allows many user level threads to be mapped to many kernel threads
- ▶ Allows the operating system to create a sufficient number of kernel threads
- ▶ Solaris prior to version 9
- ▶ Windows NT/2000 with the *ThreadFiber* package



Many-to-Many Model




Two-level Model

- ▶ Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- ▶ Examples
 - ♦ IRIX
 - ♦ HP-UX
 - ♦ Tru64 UNIX
 - ♦ Solaris 8 and earlier





Two-level Model



Thread Libraries

- ▶ **Thread library** provides programmer with API for creating and managing threads
- ▶ Two primary ways of implementing
 - ♦ Library entirely in user space
 - ♦ Kernel-level library supported by the OS



Pthreads

- ▶ May be provided either as user-level or kernel-level
- ▶ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ▶ API specifies behavior of the thread library, implementation is up to development of the library
- ▶ Common in UNIX operating systems (Solaris, Linux, Mac OS X)



Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* the thread */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```





Pthreads Example (Cont.)

```

/* get the default attributes */
pthread_attr_t attr;
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid, &attr, runner, argv[1]);
/* wait for the thread to exit */
pthread_join(tid, NULL);

printf("sum = %d\n", sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}

```

Figure 4.9 Multithreaded C program using the Pthreads API.



Win32 API Multithreaded C Program

```

#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */
/* the thread runs in this separate function */

DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;
    /* perform some basic error checking */
    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }
}

```



Win32 API Multithreaded C Program (Cont.)

```

/* create the thread */
ThreadHandle = CreateThread(
    NULL, /* default security attributes */
    0, /* default stack size */
    Summation, /* thread function */
    &Param, /* parameter to thread function */
    0, /* default creation flags */
    &ThreadId); /* returns the thread identifier */

if (ThreadHandle != NULL) {
    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
}

```



Java Threads

- ▶ Java threads are managed by the JVM
- ▶ Typically implemented using the threads model provided by underlying OS
- ▶ Java threads may be created by:
 - ✦ Extending Thread class
 - ✦ Implementing the Runnable interface





Java Multithreaded Program

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}

class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```



Java Multithreaded Program (Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>");
    }
}
```



Threading Issues

- ▶ Semantics of **fork()** and **exec()** system calls
- ▶ **Thread cancellation** of **target thread**
 - ◆ Asynchronous or deferred
- ▶ **Signal** handling
 - ◆ Synchronous and asynchronous



Threading Issues (Cont.)

- ▶ **Thread pools**
- ▶ **Thread-specific data**
 - ◆ Create Facility needed for data private to thread
- ▶ **Scheduler activations**





Semantics of fork() and exec()

- ▶ Does **fork()** duplicate only the calling thread or all threads?



Thread Cancellation

- ▶ Terminating a thread before it has finished
- ▶ Two general approaches:
 - ✦ **Asynchronous cancellation** terminates the target thread immediately.
 - ✦ **Deferred cancellation** allows the target thread to periodically check if it should be cancelled.



Signal Handling

- ▶ Signals are used in UNIX systems to notify a process that a particular event has occurred.
- ▶ A **signal handler** is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled
- ▶ Options:
 - ✦ Deliver the signal to the thread to which the signal applies
 - ✦ Deliver the signal to every thread in the process
 - ✦ Deliver the signal to certain threads in the process
 - ✦ Assign a specific thread to receive all signals for the process



Thread Pools

- ▶ Create a number of threads in a pool where they await work
- ▶ Advantages:
 - ✦ Usually slightly faster to service a request with an existing thread than create a new thread
 - ✦ Allows the number of threads in the application(s) to be bound to the size of the pool





Thread Specific Data

- ▶ Allows each thread to have its own copy of data
- ▶ Useful when you do not have control over the thread creation process (i.e., when using a thread pool)




Scheduler Activations

- ▶ Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- ▶ Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the thread library
- ▶ This communication allows an application to maintain the correct number kernel threads



Lightweight Processes



Operating System Examples

- ▶ Windows XP Threads
- ▶ Linux Thread



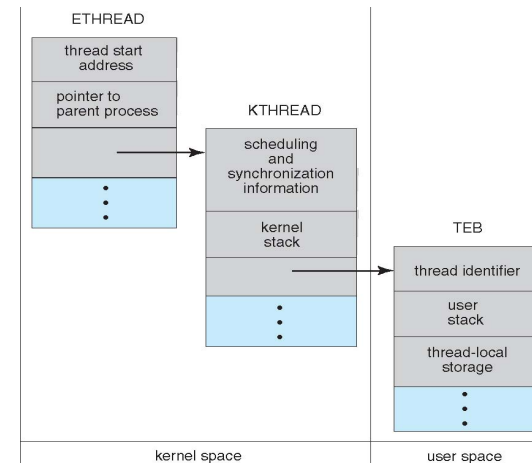


Windows XP Threads

- ▶ Implements the one-to-one mapping, kernel-level
- ▶ Each thread contains
 - ♦ A thread id
 - ♦ Register set
 - ♦ Separate user and kernel stacks
 - ♦ Private data storage area
- ▶ The register set, stacks, and private storage area are known as the **context** of the threads
- ▶ The primary data structures of a thread include:
 - ♦ ETHREAD (executive thread block)
 - ♦ KTHREAD (kernel thread block)
 - ♦ TEB (thread environment block)



Windows XP Threads Data Structures



Linux Threads

- ▶ Provides `fork()` and `clone()` system calls
 - ♦ `fork()`: traditional functionality of duplicating a process
 - ♦ `clone()`: create thread and allows a child task to share the address space of the parent task (process)
- ▶ Doesn't distinguish between process and thread; uses the term **task** – rather than **process** or **thread**
- ▶ Unique kernel data structure (`struct task_struct`) exists for each task in the system



Linux Threads

- ▶ Data of new task
 - ♦ `fork()`: new task is created; copy of all the associated data structures of the parent process
 - ♦ `clone()`: new task is created and **points** to the data structures of the parent task, depending on the set of flags passed to `clone()`
 - ♦ if none of these flags is set: no sharing takes place, functionality similar to that provided by the `fork()` system call

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.



End of Chapter 4

