



Secure Architecture Principles

- Isolation and Least Privilege
- Access Control Concepts
- Operating Systems
- Browser Isolation and Least Privilege



Secure Architecture Principles

Isolation and Least Privilege


Principles of Secure Design

- Compartmentalization
 - Isolation
 - Principle of least privilege
- Defense in depth
 - Use more than one security mechanism
 - Secure the weakest link
 - Fail securely
- Keep it simple


Principle of Least Privilege

- Principle of Least Privilege
 - A system module should only have the minimal privileges needed for its intended purposes
- What's a privilege?
 - Ability to access or modify a resource
- Assumes compartmentalization and isolation
 - Separate the system into isolated compartments
 - Limit interaction between compartments

Monolithic design




Monolithic design




Monolithic design




Component design



Component design



Component design



Principle of Least Privilege

- Principle of Least Privilege
 - A system module should only have the minimal privileges needed for its intended purposes
- What's a privilege?
 - Ability to access or modify a resource
- Assumes compartmentalization and isolation
 - Separate the system into isolated compartments
 - Limit interaction between compartments

Example: Mail Agent

- Requirements
 - Receive and send email over external network
 - Place incoming email into local user inbox files
- Sendmail
 - Traditional Unix
 - Monolithic design
 - Historical source of many vulnerabilities
- Qmail
 - Compartmentalized design

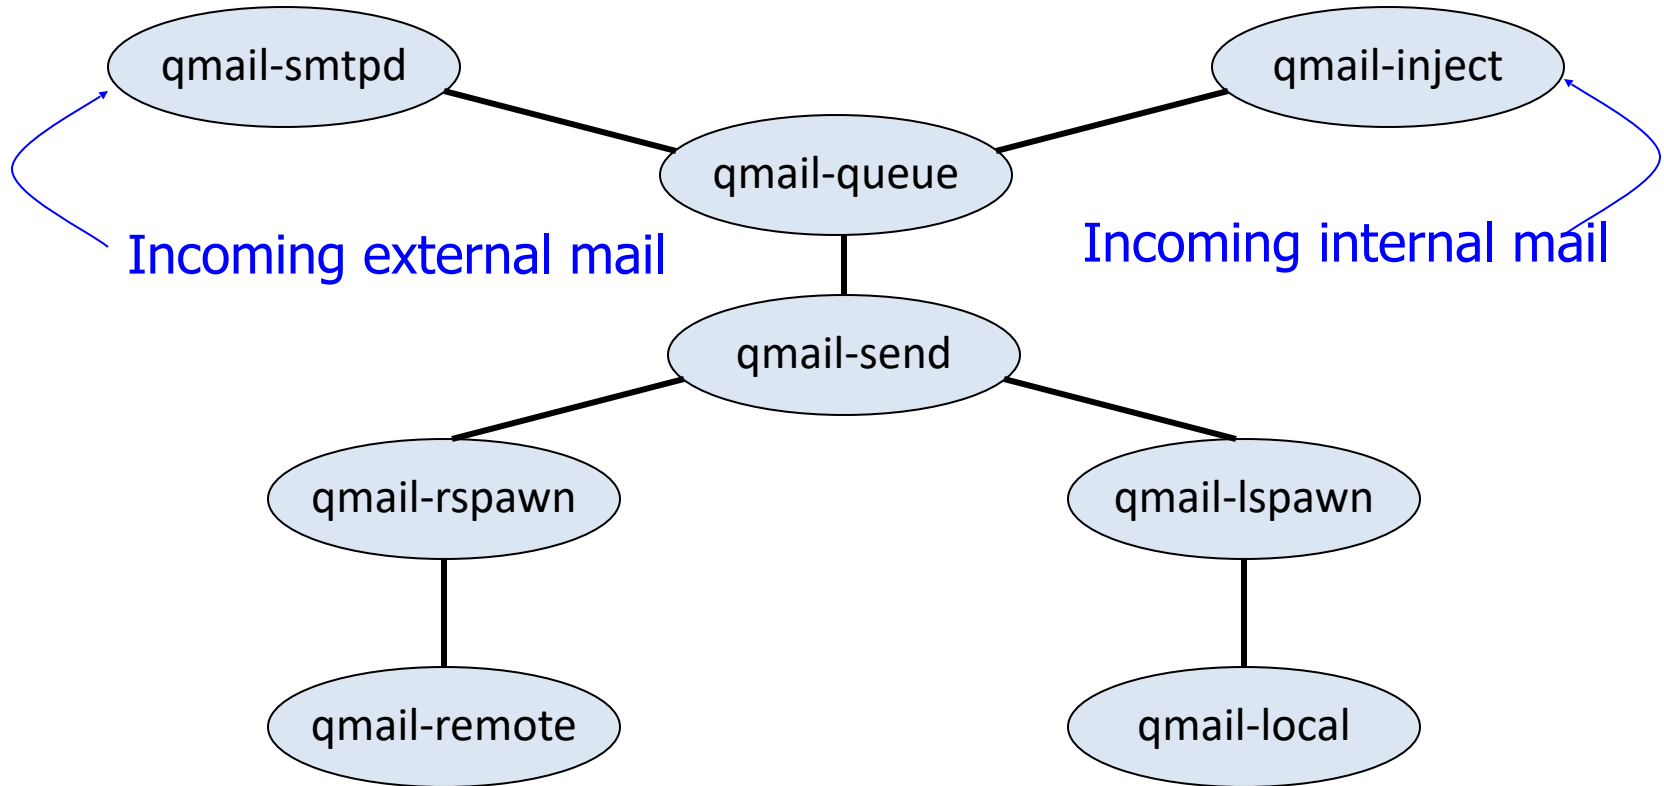
OS Basics (before examples)

- Isolation between processes
 - Each process has a UID
 - Two processes with same UID have same permissions
 - A process may access files, network sockets,
 - Permission granted according to UID
- Relation to previous terminology
 - Compartment defined by UID
 - Privileges defined by actions allowed on system resources

Qmail design

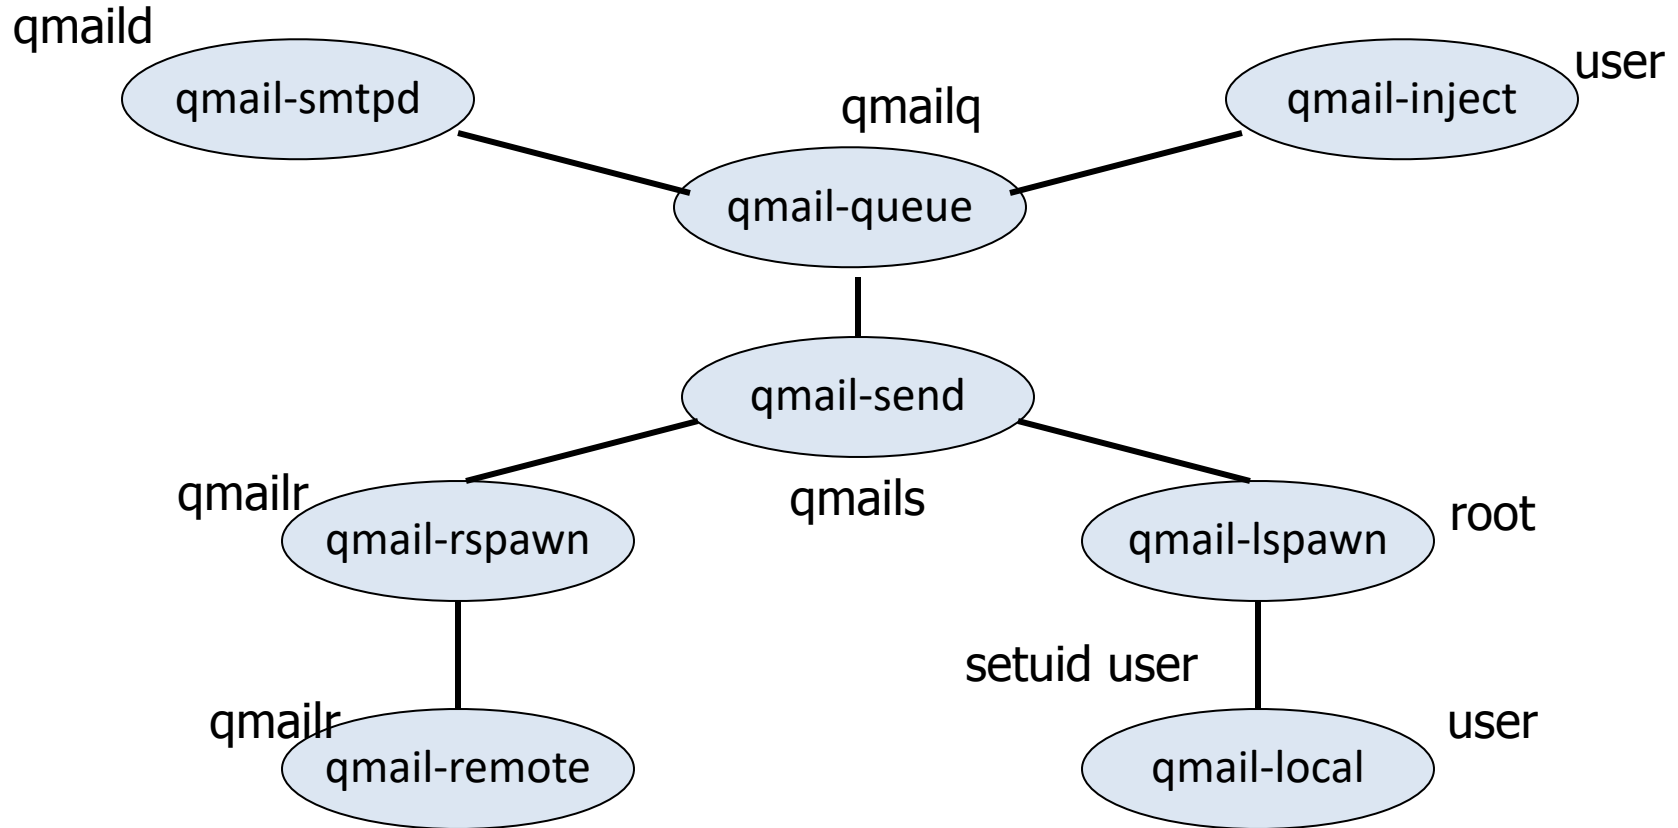
- Isolation based on OS isolation
 - Separate modules run as separate “users”
 - Each user only has access to specific resources
- Least privilege
 - Minimal privileges for each UID
 - Only one “setuid” program
 - setuid allows a program to run as different users
 - Only one “root” program
 - root program has all privileges

Structure of qmail

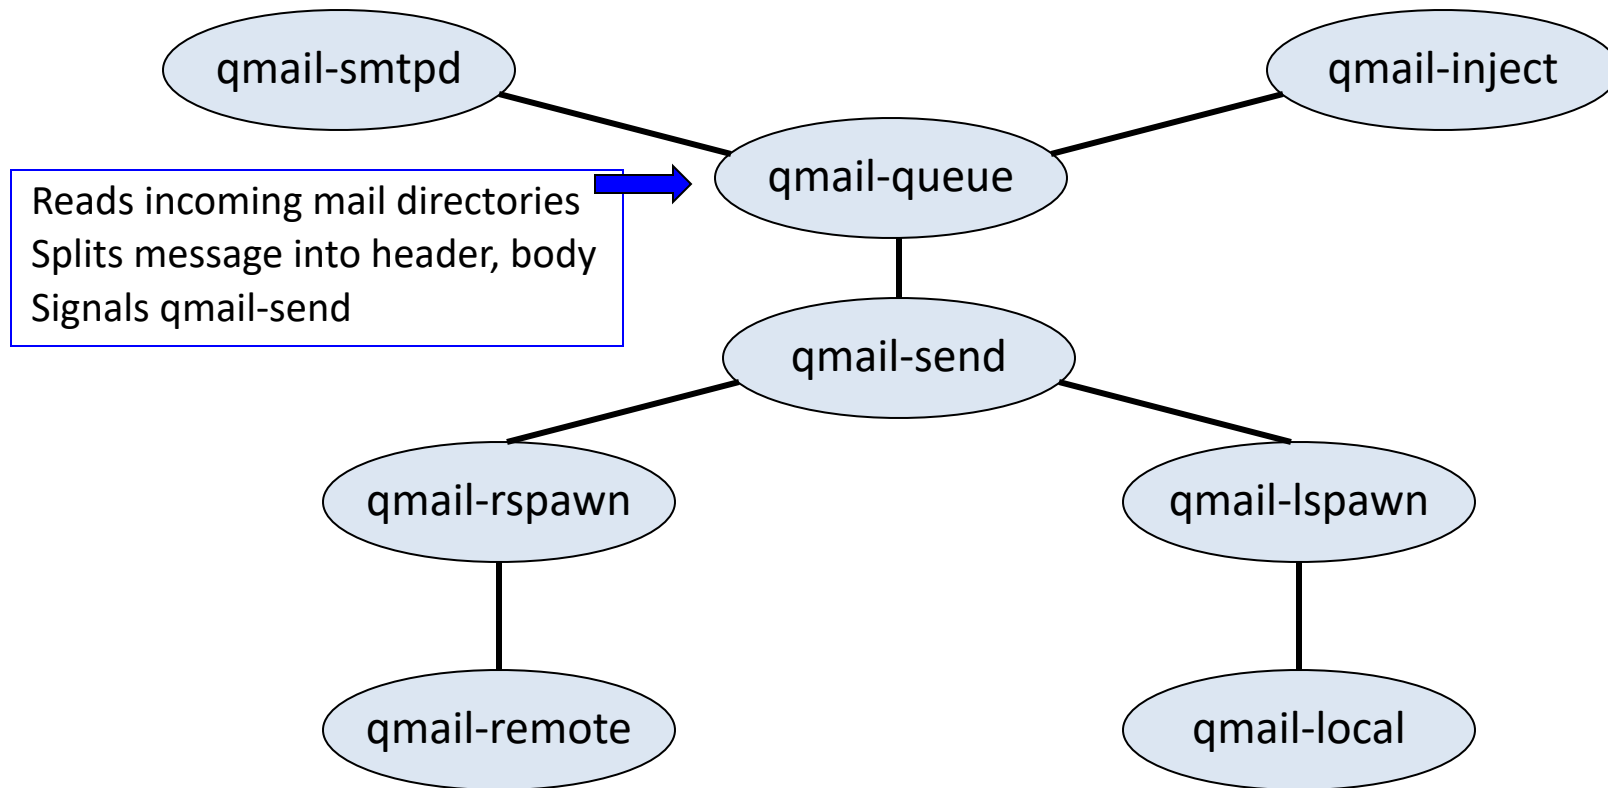


Isolation by Unix UIDs

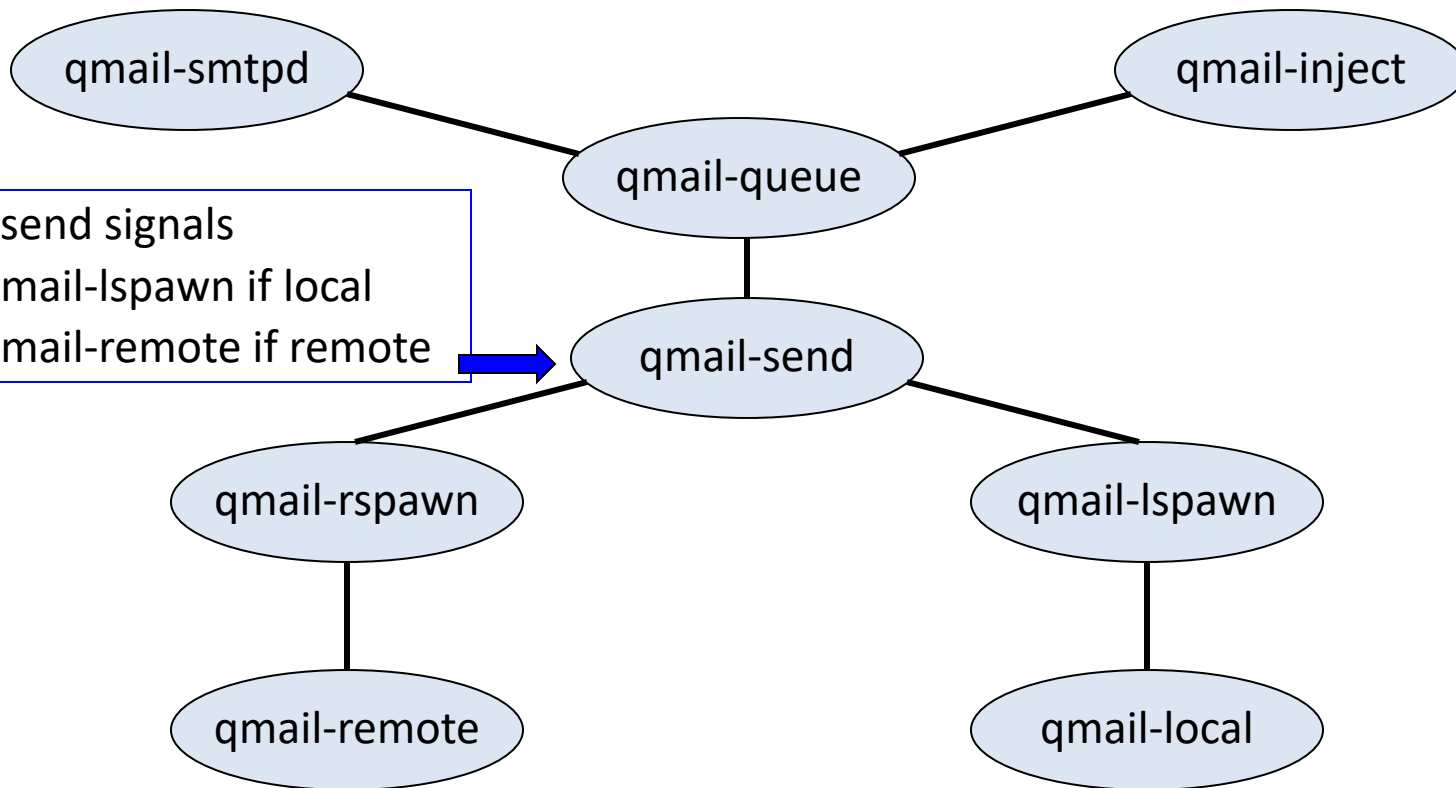
qmailq – user who is allowed to read/write mail queue



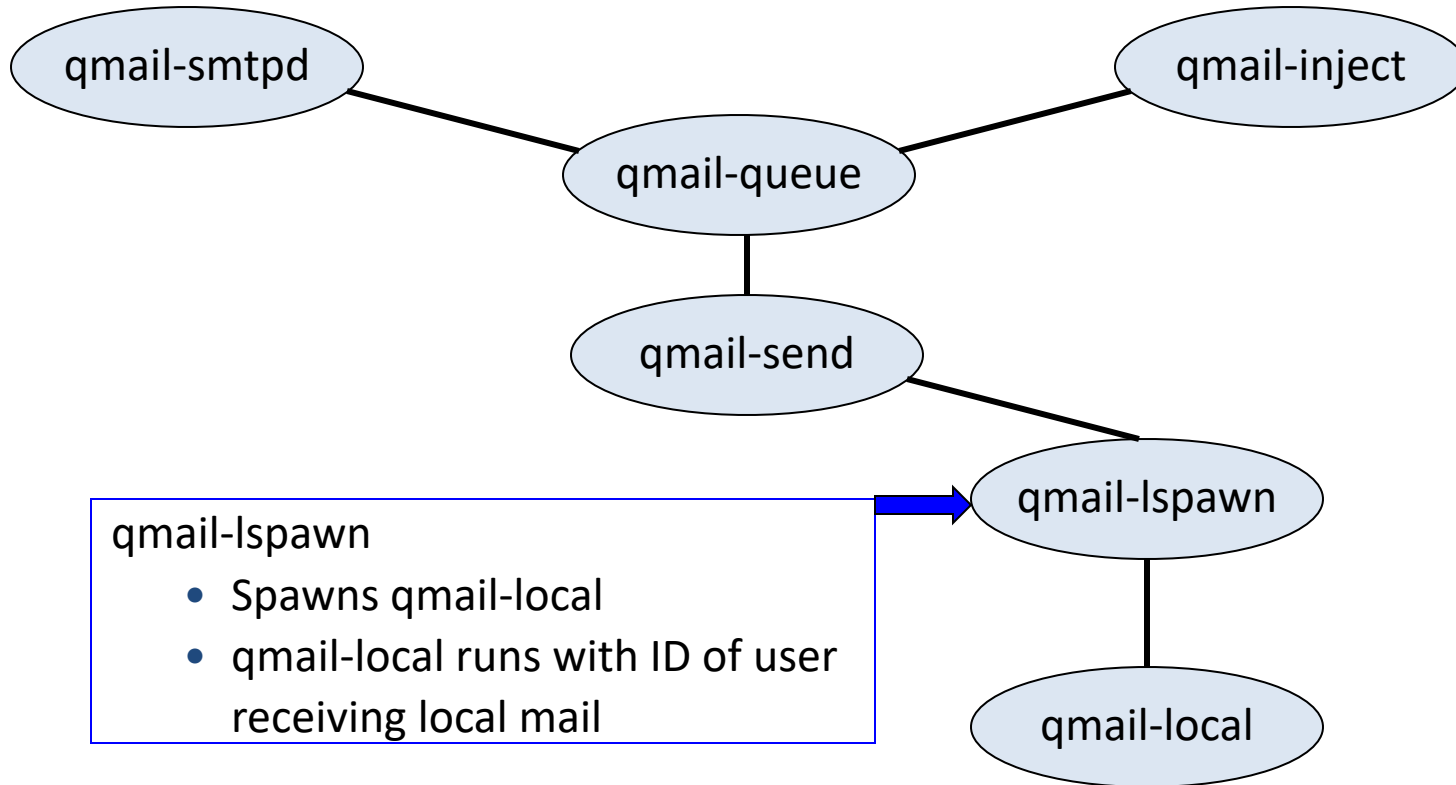
Structure of qmail




Structure of qmail



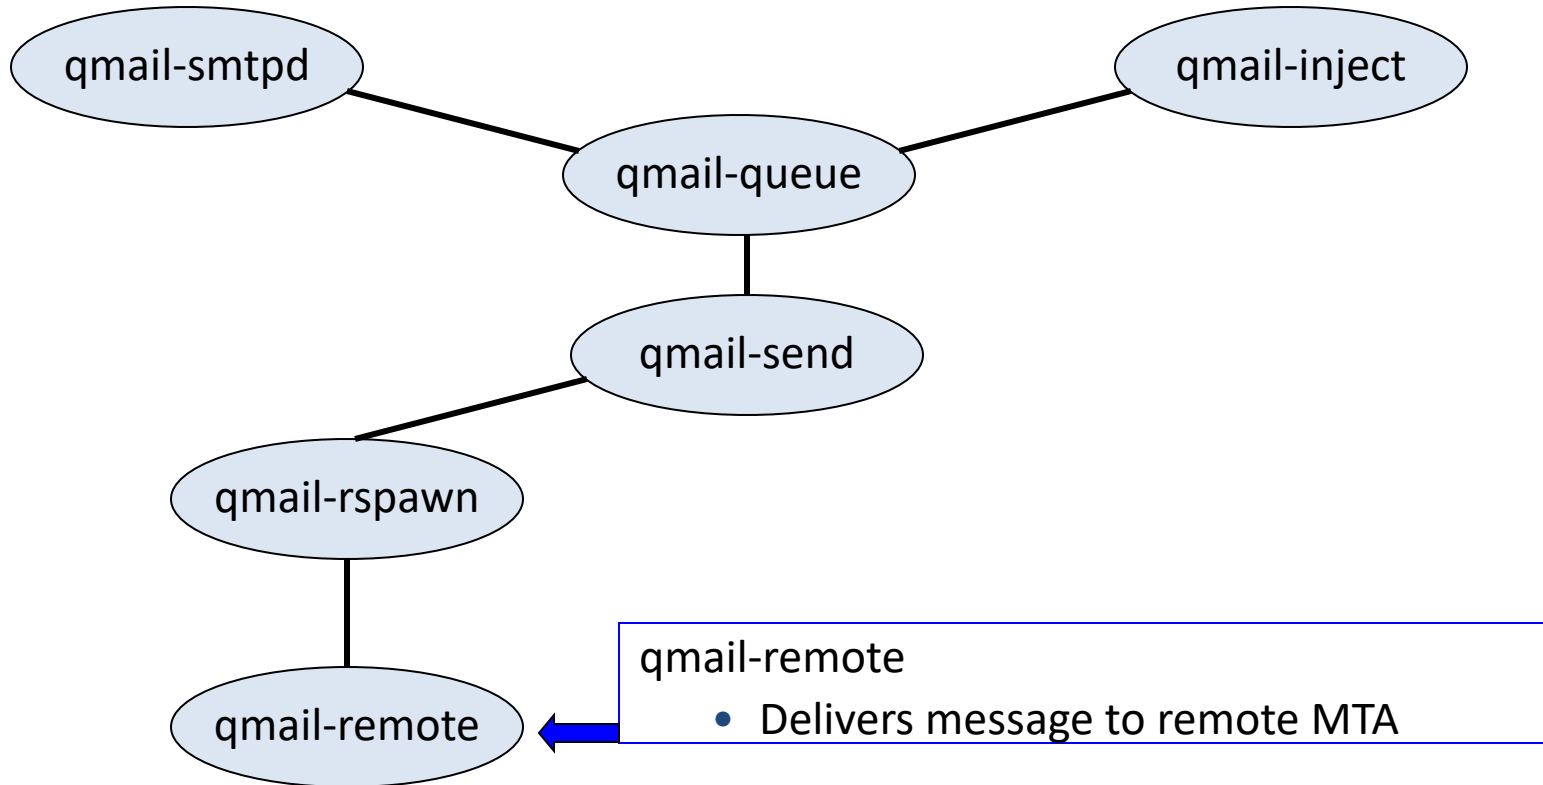
Structure of qmail



Structure of qmail

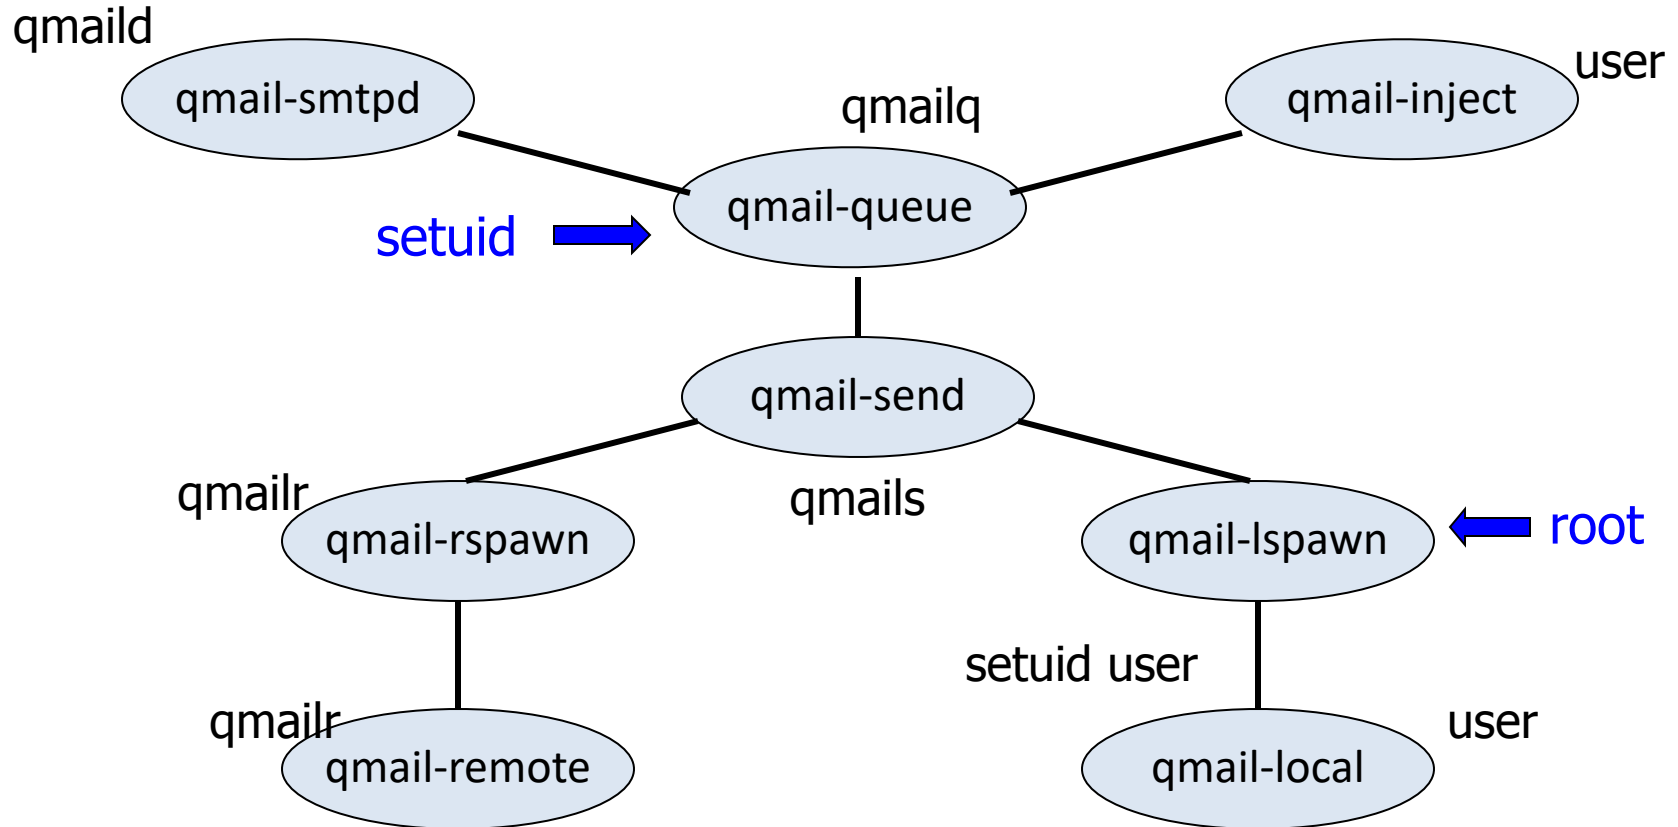


Structure of qmail




Isolation by Unix UIDs

qmailq – user who is allowed to read/write mail queue



Least privilege

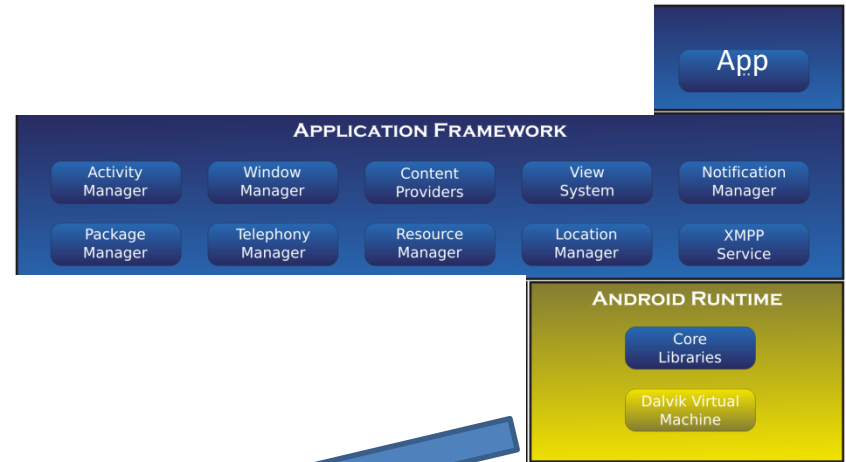
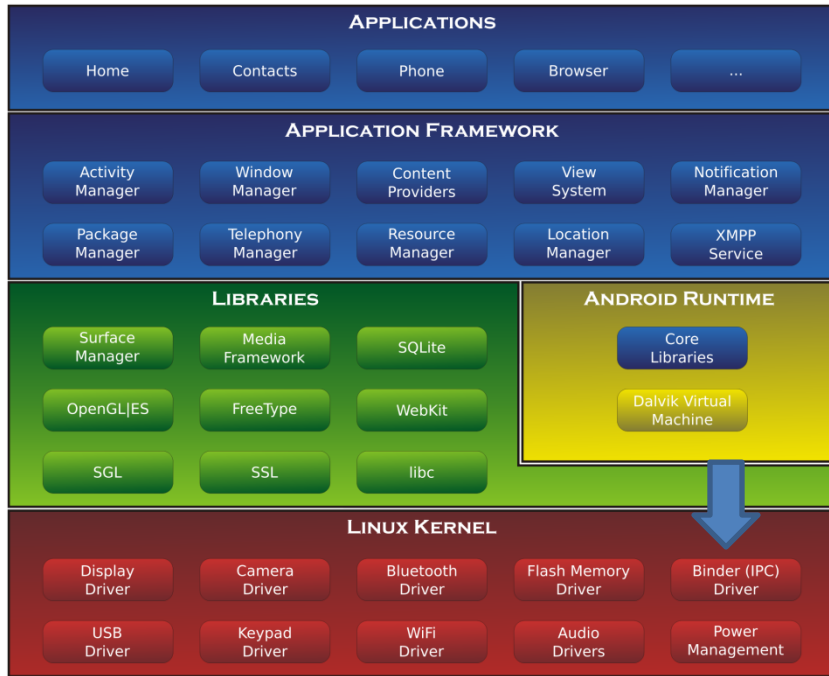


Android process isolation

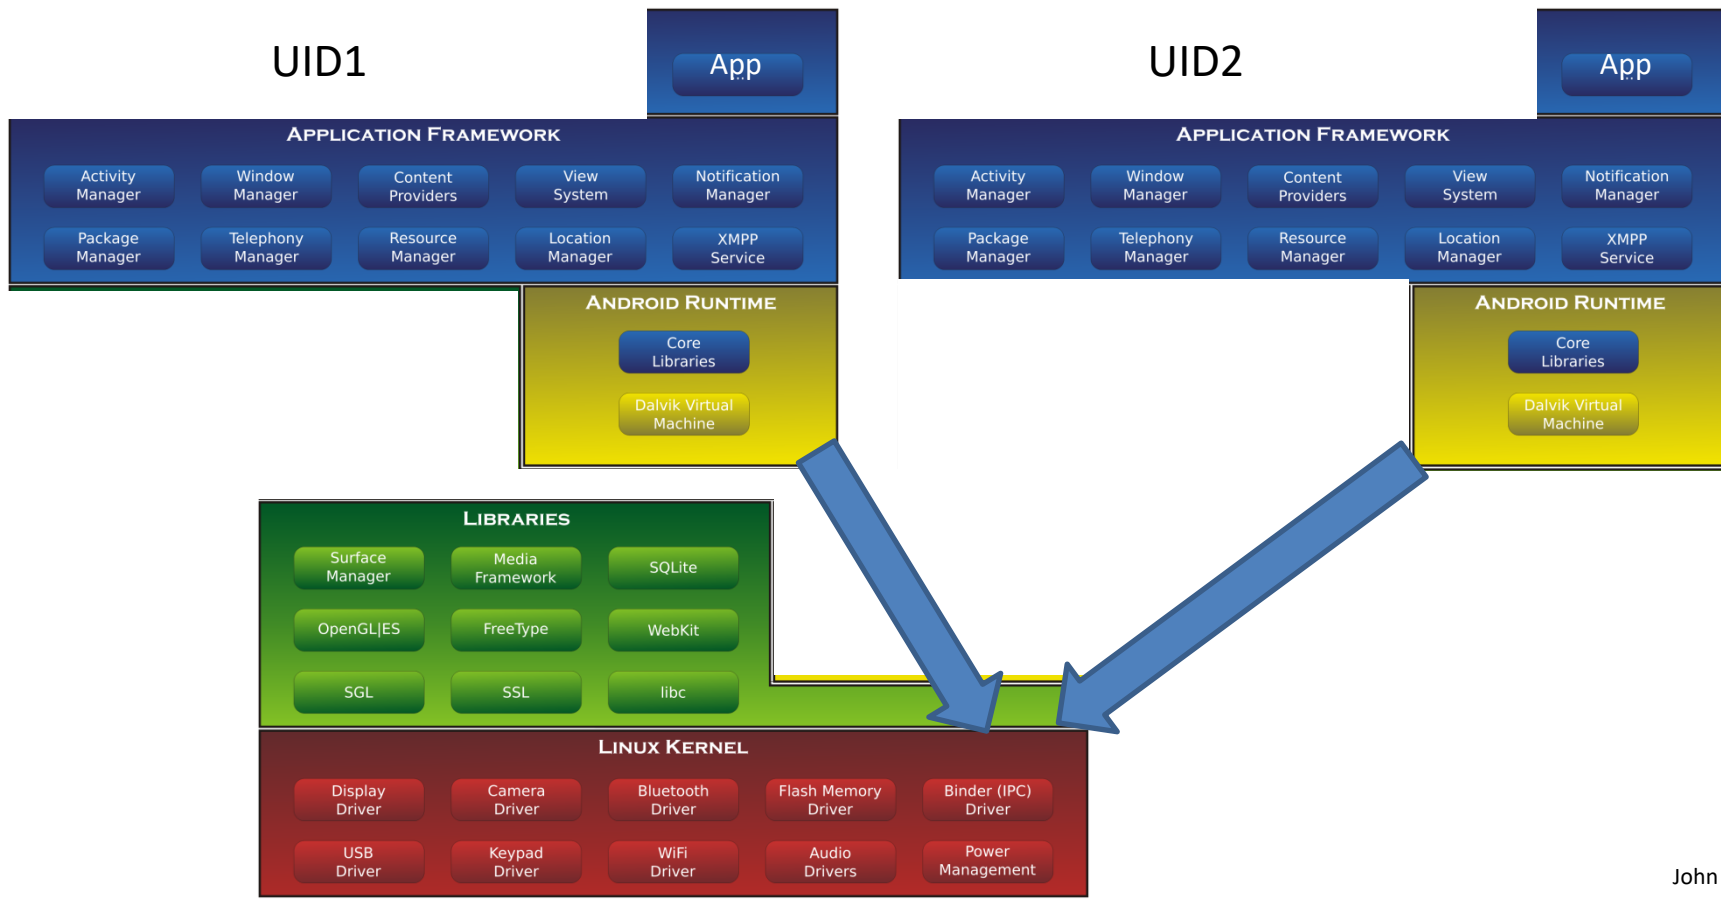
- Android application sandbox
 - Isolation: Each application runs with its own UID in own VM
 - Provides memory protection
 - Communication limited to using Unix domain sockets
 - Only ping, zygote (spawn another process) run as root
 - Interaction: reference monitor checks permissions on inter-component communication
 - Least Privilege: Applications announces permission
 - User grants access at install time



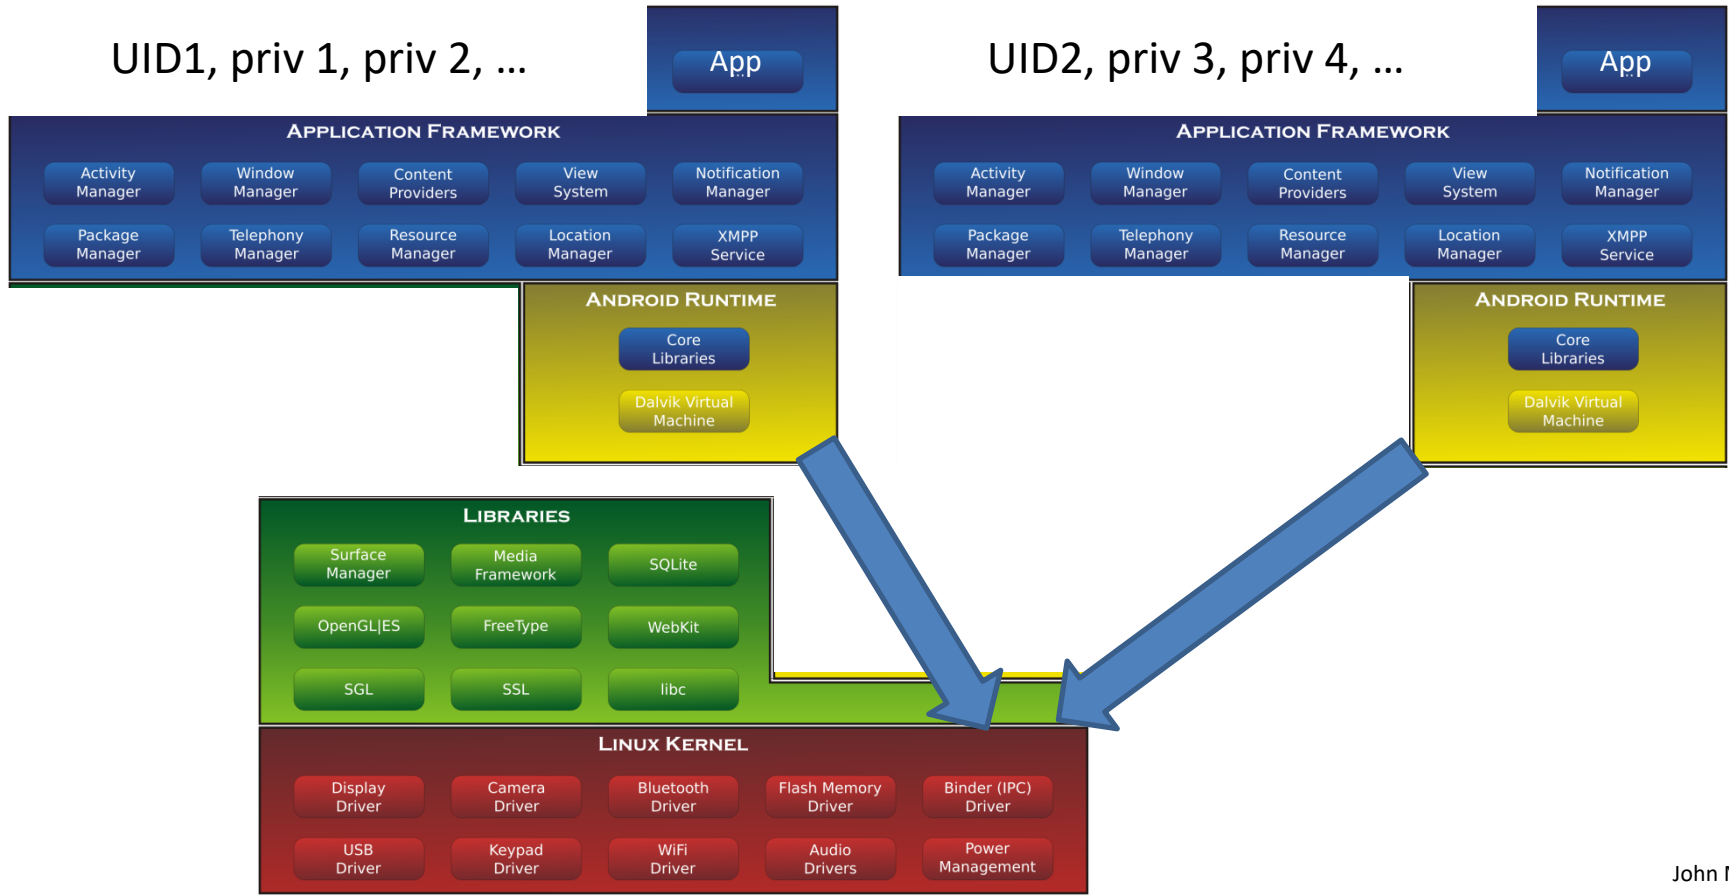
Isolation: different apps under different UIDs



Isolation: different apps under different UIDs



Privileges set at install time



Discussion?

- Principle of Least Privilege
- Qmail example
- Android app sandbox example




Secure Architecture Principles

Access Control Concepts

Access control

- Assumptions
 - System knows who the user is
 - Authentication via name and password, other credential
 - Access requests pass through gatekeeper (reference monitor)
 - System must not allow monitor to be bypassed



Access control matrix [Lampson]

The diagram illustrates an Access Control Matrix. A large curly bracket on the left side groups the rows under the label "Subjects". A large curly bracket above the columns groups them under the label "Objects". The matrix is a table with 6 columns and 6 rows. The first column is for Subjects, and the other five columns are for Objects. The cells contain access permissions: "read" or "write" for granted permissions, and "-" for denied permissions.

	File 1	File 2	File 3	...	File n
User 1	read	write	-	-	read
User 2	write	write	write	-	-
User 3	-	-	-	read	read
...					
User m	read	write	read	write	read

Implementation concepts

- Access control list (ACL)
 - Store column of matrix with the resource
- Capability
 - User holds a “ticket” for each resource
 - Two variations
 - store row of matrix with user, under OS control
 - unforgeable ticket in user space

	File 1	File 2	...
User 1	read	write	-
User 2	write	write	-
User 3	-	-	read
...			
User m	Read	write	write

Access control lists are widely used, often with groups

Some aspects of capability concept are used in many systems

ACL: my name is on the list



Capability: I have a ticket



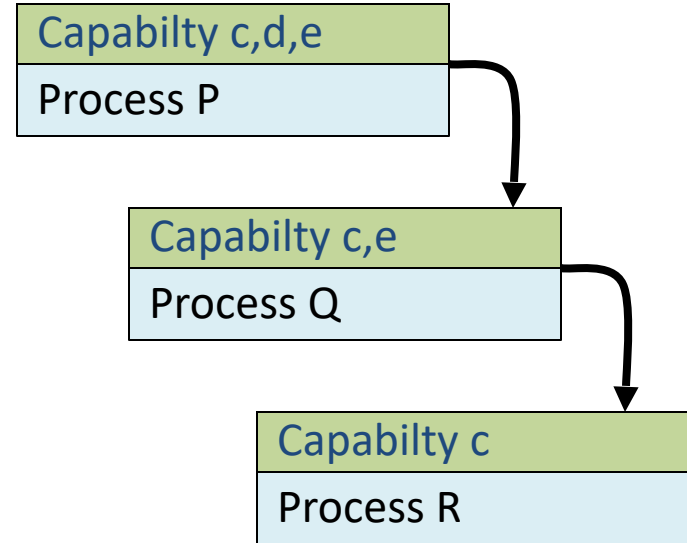
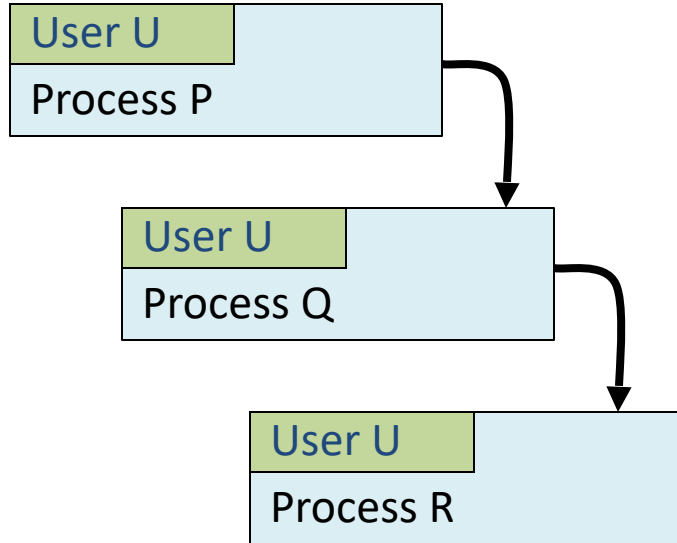
ACL vs Capabilities

- Access control list
 - Associate list with each object
 - Check user/group against list
 - Relies on authentication: need to know user
- Capabilities
 - Capability is unforgeable ticket
 - Random bit sequence (or managed by OS)
 - Can be passed from one process to another
 - Reference monitor checks ticket
 - Does not need to know identify of user/process

ACL vs Capabilities


- Delegation
 - Cap: Process can pass capability at run time
 - ACL: Try to get owner to add permission to list?
 - More common: let other process act under current user
- Revocation
 - ACL: Remove user or group from list
 - Cap: Try to get capability back from process?
 - Possible in some systems if appropriate bookkeeping
 - OS knows which data is capability
 - If capability is used for multiple resources, have to revoke all or none ...
 - Indirection: capability points to pointer to resource
 - If $C \rightarrow P \rightarrow R$, then revoke capability C by setting $P=0$

Process creation: ACL vs Capabilities




Roles (aka Groups)

- Role = set of users
 - Administrator, PowerUser, User, Guest
 - Assign permissions to roles; each user gets permission
- Role hierarchy
 - Partial order of roles
 - Each role gets permissions of roles below
 - List only new permissions given to each role



Role-Based Access Control



Advantage: users change more frequently than roles

Access control summary

- Access control involves reference monitor
 - Check permissions: $\langle \text{user info, action} \rangle \rightarrow \text{yes/no}$
 - Important: must be no way to bypass this check
- Access control matrix
 - Two implementations: access control lists vs capabilities
 - Advantages and disadvantages of each
- Role-based access control
 - Use group as “user info”; use group hierarchies

Discussion?

- Access control matrix
 - What are the advantages of access control lists (ACL)
 - What are the advantages of capabilities
- Role-based access control
 - Why is this helpful?




Secure Architecture Principles

Operating Systems

Unix

- What access control concepts are used?
 - Truncated access control list
 - A form of role-based access control

	File 1	File 2	...
User 1	read	write	-
User 2	write	write	-
User 3	-	-	read
...			
Role r	Read	write	write



	File 1	File 2	...
Owner	read	write	-
Group	write	write	-
Other	-	-	read

Unix access control

- Process has user id
 - Inherit from creating process
 - Process can change id
 - Restricted set of options
 - Special “root” id
 - All access allowed
- File has access control list (ACL)
 - Grants permission to users
 - Three “roles”: owner, group, other

	File 1	File 2	...
Owner	read	write	-
Group	write	write	-
Other	-	-	read

Unix file access control list

- Each file has owner and group
- Permissions set by owner
 - Read, write, execute
 - Owner, group, other
 - Represented by vector of four octal values
- Only owner, root can change permissions
 - This privilege cannot be delegated or shared
- Setid bits – Discuss in a few slides

The diagram illustrates the structure of Unix file permissions. It consists of three groups of permissions, each represented by a blue bracket above a label. The first group is labeled 'ownr' and contains the permissions 'rwx'. The second group is labeled 'grp' and contains the permissions 'rwx'. The third group is labeled 'othr' and contains the permissions 'rwx'.

Category	Permissions
ownr	rwx
grp	rwx
othr	rwx

Example directory listing

<i>access</i>	<i>owner</i>	<i>group</i>	<i>size</i>	<i>modification</i>	<i>name</i>
-rw-rw-r--	pbg	staff	31200	Sep 3 08:30	intro.ps
drwx-----	pbg	staff	512	Jul 8 09:33	private/
drwxrwxr-x	pbg	staff	512	Jul 8 09:35	doc/
drwxrwx---	jwg	student	512	Aug 3 14:13	student-proj/
-rw-r--r--	pbg	staff	9423	Feb 24 2012	program.c
-rwxr-xr-x	pbg	staff	20471	Feb 24 2012	program
drwx--x--x	tag	faculty	512	Jul 31 10:31	lib/
drwx-----	pbg	staff	1024	Aug 29 06:52	mail/
drwxrwxrwx	pbg	staff	512	Jul 8 09:35	test/

Process effective user id (EUID)

- Each process has three Ids (+ more under Linux)
 - Real user ID (RUID)
 - same as the user ID of parent (unless changed)
 - used to determine which user started the process
 - Effective user ID (EUID)
 - from set user ID bit on the file being executed, or sys call
 - determines the permissions for process
 - file access and port binding
 - Saved user ID (SUID)
 - So previous EUID can be restored
- Real group ID, effective group ID, used similarly

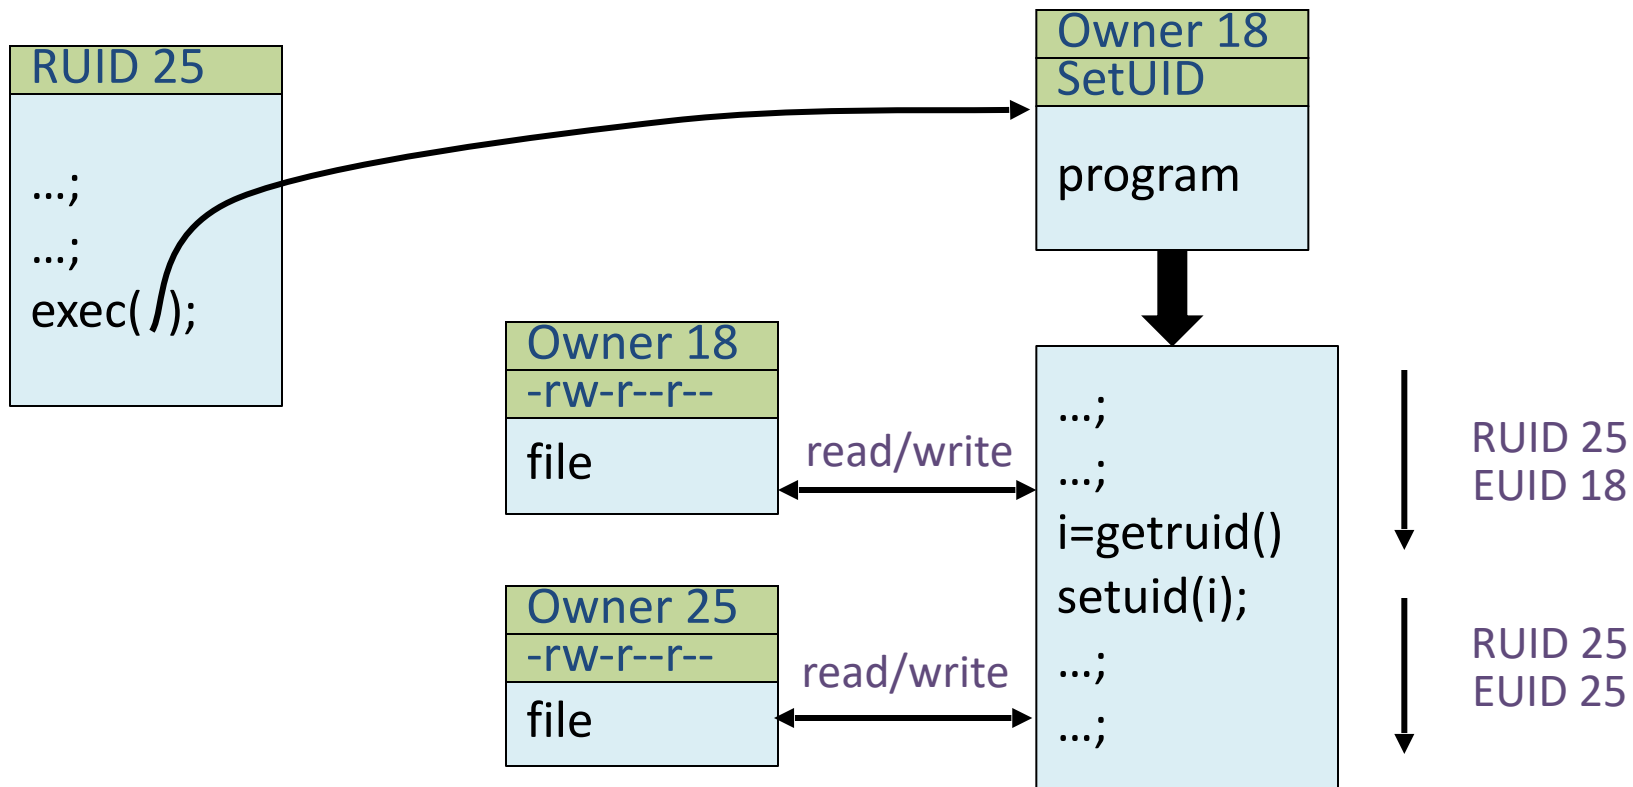
Process Operations and IDs

- Root
 - ID=0 for superuser root; can access any file
- Fork and Exec
 - Inherit three IDs, except exec of file with setuid bit
- Setuid system call
 - seteuid(newid) can set EUID to
 - Real ID or saved ID, regardless of current EUID
 - Any ID, if EUID is root
- Details are actually more complicated
 - Several different calls: setuid, seteuid, setreuid

Setid bits on executable Unix file

- Three setid bits
 - Setuid – set EUID of process to ID of file owner
 - Setgid – set EGID of process to GID of file
 - Sticky
 - Off: if user has write permission on directory, can rename or remove files, even if not owner
 - On: only file owner, directory owner, and root can rename or remove file in the directory

Example



Unix access control summary

- Good things
 - Some protection from most users
 - Flexible enough to make practical systems possible
- Main limitation
 - Coarse-grained ACLs – user, group, other
 - Too tempting to use root privileges
 - No way to assume some root privileges without all

Weakness in unix isolation, privileges

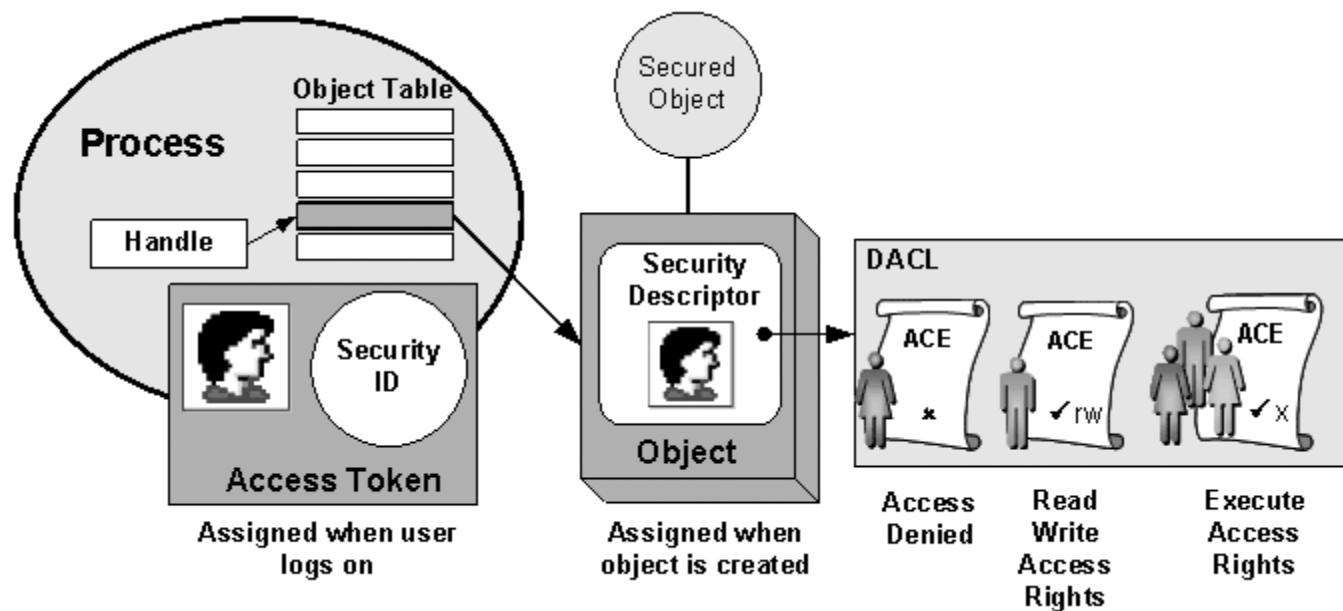
- Network-facing Daemons
 - Root processes with network ports open to all remote parties, e.g., sshd, ftpd, sendmail, ...
- Rootkits
 - System extension via dynamically loaded kernel modules
- Environment Variables
 - System variables such as LIBPATH that are shared state across applications. An attacker can change LIBPATH to load an attacker-provided file as a dynamic library

Weakness in unix isolation, privileges

- Shared Resources
 - Since any process can create files in /tmp directory, an untrusted process may create files that are used by arbitrary system processes
- Time-of-Check-to-Time-of-Use (TOCTTOU)
 - Typically, a root process uses system call to determine if initiating user has permission to a particular file, e.g. /tmp/X.
 - After access is authorized and before the file open, user may change the file /tmp/X to a symbolic link to a target file /etc/shadow.

Access control in Windows

- Full access control lists
 - Specify access for groups and users
 - Read, modify, change owner, delete
- Some additional concepts
 - Tokens
 - Security attributes
- Generally, more precise, more flexible than Unix
 - Can define new permissions
 - Can transfer some but not all privileges (*cf.* capabilities)



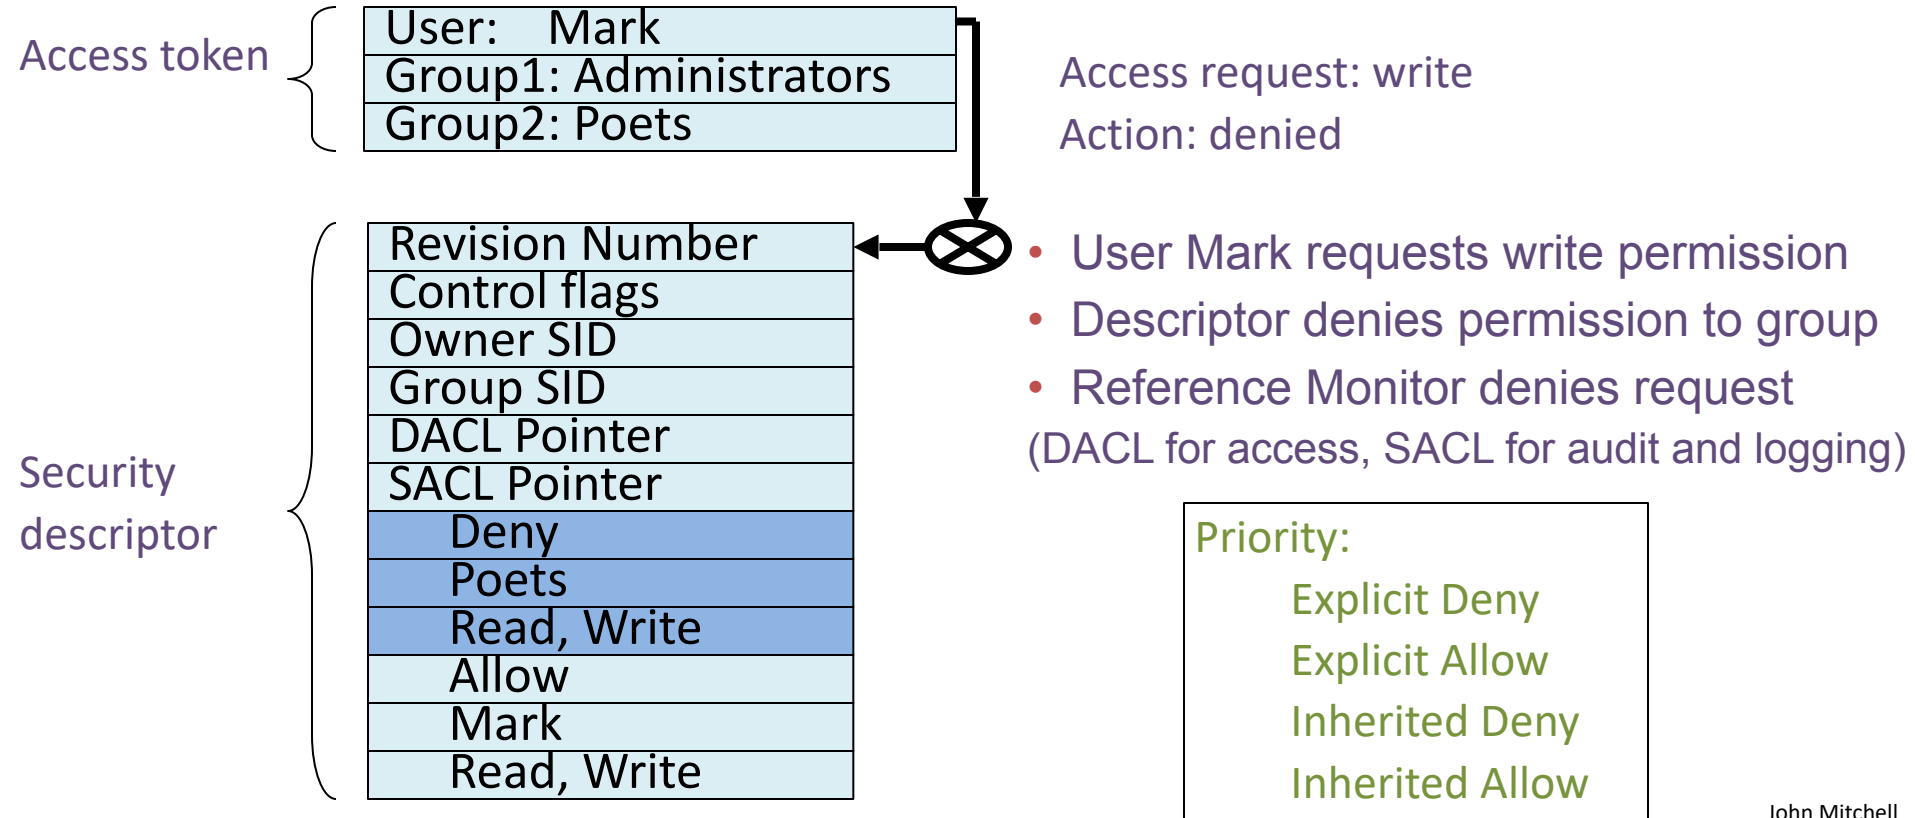
Process has set of tokens

- Called the process “security context”
 - Privileges, accounts, and groups associated with the process or thread
 - Presented as set of tokens
- Interesting feature: impersonation token
 - Used temporarily to adopt a different security context, usually of another user (similar to use of capability/setuid)

Object has security descriptor

- Specifies who can perform what actions on the object
 - Header (revision number, control flags, ...)
 - SID of the object's owner
 - SID of the primary group of the object
 - Two attached optional lists:
 - Discretionary Access Control List (DACL) – users, groups, ...
 - System Access Control List (SACL) – system logs, ..

Example access request



Impersonation Tokens (compare to setuid)

- Process adopts security attributes of another
 - Client passes impersonation token to server
- Client specifies impersonation level
 - Anonymous
 - Token has no information about the client
 - Identification
 - Obtain the SIDs of client and client's privileges, but server cannot impersonate the client
 - Impersonation
 - Impersonate the client
 - Delegation
 - Lets server impersonate client on local, remote systems

Windows access control summary

- Full access control lists
 - Specify access for groups and users
 - Read, modify, change owner, delete
- Some additional concepts
 - Tokens
 - Security attributes
- Generally, more precise, more flexible than Unix
 - Can define new permissions
 - Can transfer some but not all privileges (*cf.* capabilities)

Weakness in isolation, privileges

- Similar problems to Unix
 - E.g., Rootkits leveraging dynamically loaded kernel modules
- Windows Registry
 - Global hierarchical database to store data for all programs
 - Registry entry can be associated with a security context that limits access; common to be able to write sensitive entry
- Can have permissions enabled by default
 - Historically, many Windows deployments also came with full permissions and functionality enabled

Discussion?

- Unix access control
 - What information is associated with a process?
 - What information is associated with a resource (file)?
 - How are they compared?
 - What form of delegation of authority is possible?
- Windows access control
 - What information is associated with a process?
 - What information is associated with a resource (file)?
 - How are they compared?
 - What form of delegation of authority is possible?
- Comparison, pros and cons?



Secure Architecture Principles

Browser Isolation and Least Privilege

Let's look at browser example

- Browser is an execution environment
 - Has access control policies similar to an OS
- Browser runs under control of an OS
 - Use least privilege to keep the browser code secure against attacks that would break the browser enforcement of web security policy

Topic here: implementation of browser using least privilege

Web browser: an analogy

Operating system

- Subject: Processes
 - Has User ID (UID, SID)
 - Discretionary access control
- Objects
 - File
 - Network
 - ...
- Vulnerabilities
 - Untrusted programs
 - Buffer overflow
 - ...

Web browser

- Subject: web content (JavaScript)
 - Has “Origin”
 - Mandatory access control
- Objects
 - Document object model
 - Frames
 - Cookies / localStorage
- Vulnerabilities
 - Cross-site scripting
 - Implementation bugs
 - ...


The web browser enforces its own internal policy. If the browser implementation is corrupted, this mechanism becomes unreliable.

Components of security policy

- Frame-Frame relationships
 - `canScript(A,B)`
 - Can Frame A execute a script that manipulates arbitrary/nontrivial DOM elements of Frame B?
 - `canNavigate(A,B)`
 - Can Frame A change the origin of content for Frame B?
- Frame-principal relationships
 - `readCookie(A,S)`, `writeCookie(A,S)`
 - Can Frame A read/write cookies from site S?

Chromium Security Architecture

- Browser ("kernel")
 - Full privileges (file system, networking)
- Rendering engine
 - Can have multiple processes
 - Sandboxed
- One process per plugin
 - Full privileges of browser

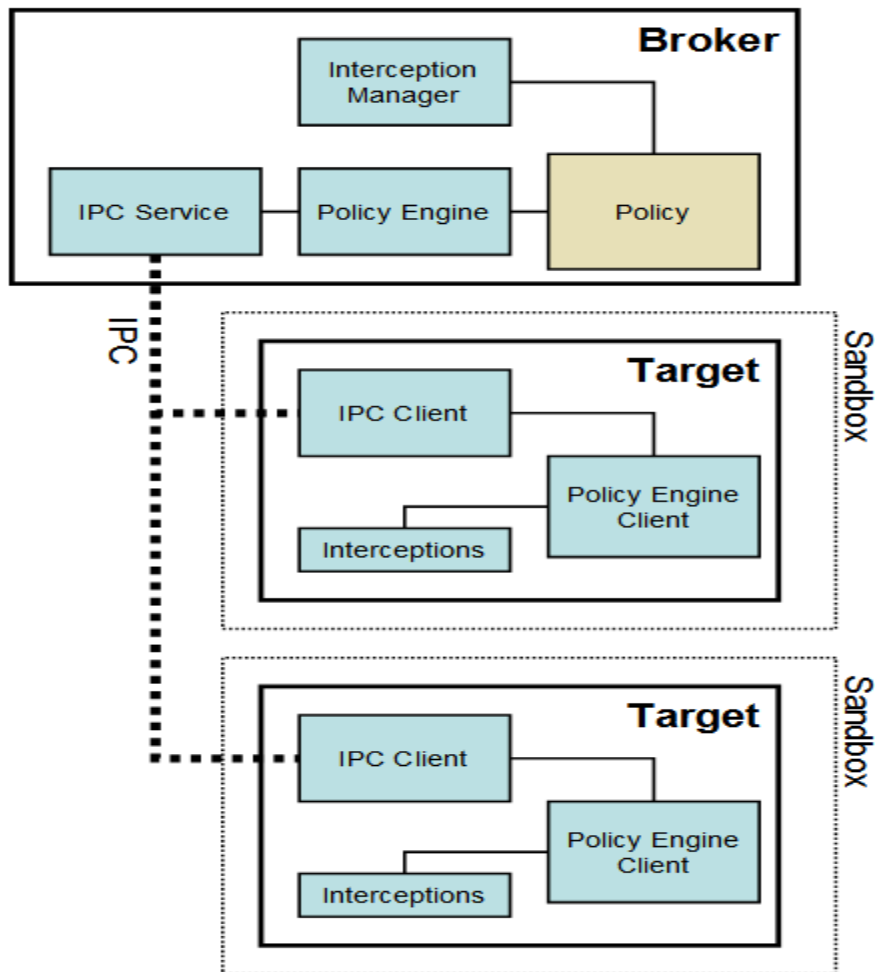


Task Allocation

Rendering Engine	Browser Kernel
HTML parsing	Cookie database
CSS parsing	History database
Image decoding	Password database
JavaScript interpreter	Window management
Regular expressions	Location bar
Layout	Safe Browsing blacklist
Document Object Model	Network stack
Rendering	SSL/TLS
SVG	Disk cache
XML parsing	Download manager
XSLT	Clipboard
Both	
URL parsing	
Unicode parsing	

Chromium

Communicating sandbox components



See: <http://dev.chromium.org/developers/design-documents/sandbox/>

Leverage OS Isolation

- Sandbox based on four OS mechanisms (e.g., Windows)
 - A restricted token
 - The Windows *job* object
 - The Windows *desktop* object
 - Windows *integrity levels*
- Specifically, the rendering engine
 - adjusts security token by converting SIDS to DENY_ONLY, adding restricted SID, and calling AdjustTokenPrivileges
 - runs in a Windows Job Object, restricting ability to create new processes, read or write clipboard, ..
 - runs on a separate desktop, mitigating lax security checking of some Windows APIs

See: <http://dev.chromium.org/developers/design-documents/sandbox/>

Evaluation: CVE count

- Total CVEs:

	Browser	Renderer	Unclassified
Internet Explorer	4	10	5
Firefox	17	40	3
Safari	12	37	1

- Arbitrary code execution vulnerabilities:

	Browser	Renderer	Unclassified
Internet Explorer	1	9	5
Firefox	5	19	0
Safari	5	10	0

Discussion?

- How does Chrome architecture use principle of least privilege?
 - What are the isolated modules?
 - Which privileges are given to each module?
- Why is this effective?
- Are there other ways you could use operating system features to improve isolation and least privilege?

Summary

- Security principles
 - Isolation
 - Principle of Least Privilege
 - Qmail, Android examples
- Access Control Concepts
 - Matrix, ACL, Capabilities
- OS Mechanisms
 - Unix: UID, ACL, Setuid
 - Windows: SID, Tokens, Security Descriptor, Impersonation
- Browser security architecture
 - Isolation and least privilege example