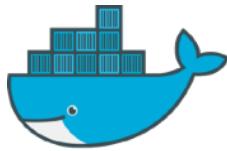


CS155: Infrastructure Security Building Blocks



@diogomonica





- Infrastructure Company
- Running on 10% of the internet
- Infra: Go, some python



- Mobile payments company
- Moving \$60 billion annually
- Infra: Java & Ruby, some Go

Agenda

- 1. Why infrastructure security?**
- 2. Biased timeline of how we got here**
- 3. What is Docker**
- 4. Docker security tetrodinos**
- 5. Demo**

Why Infrastructure Security?

~1980

2000



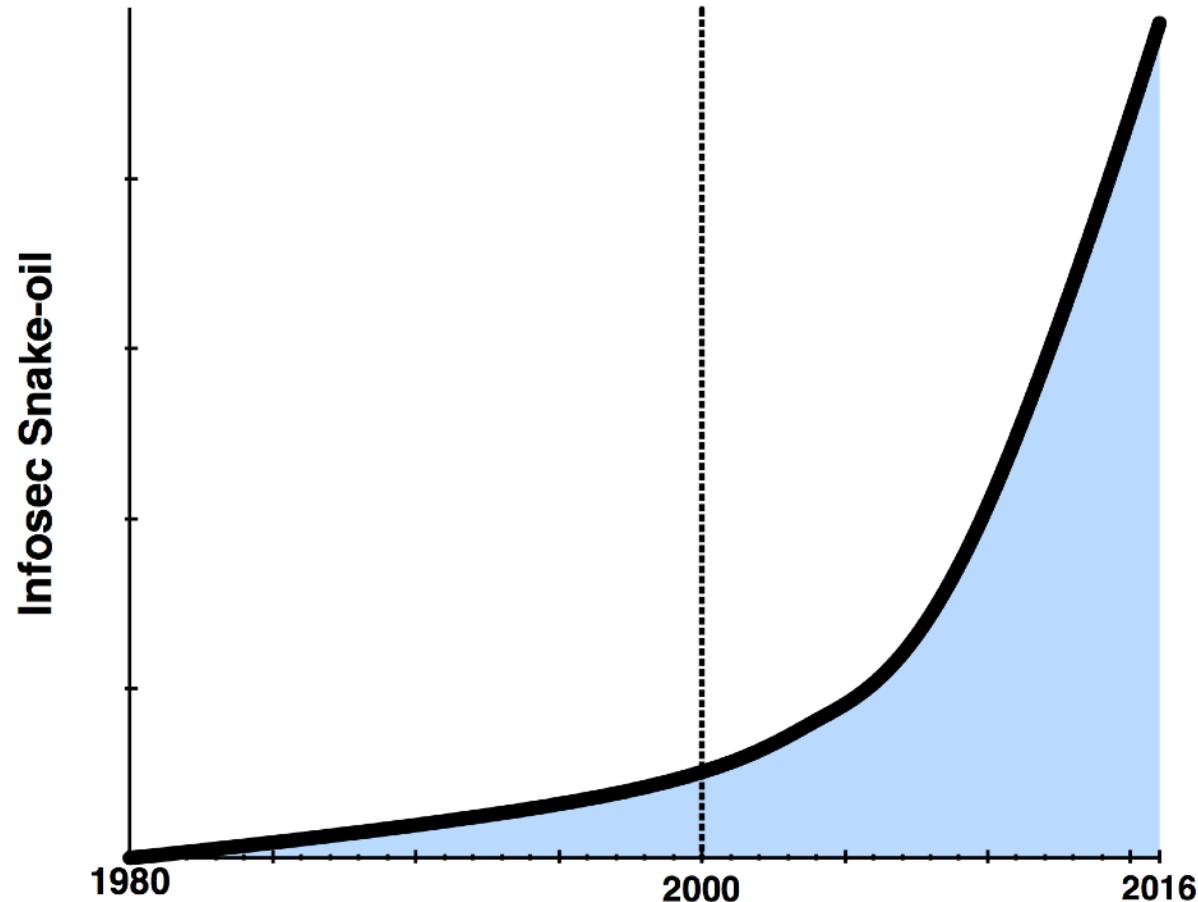
The Golden Ages

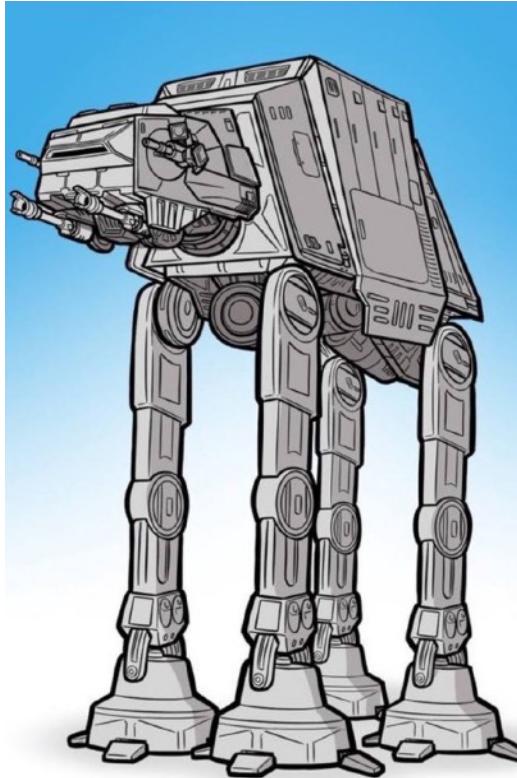


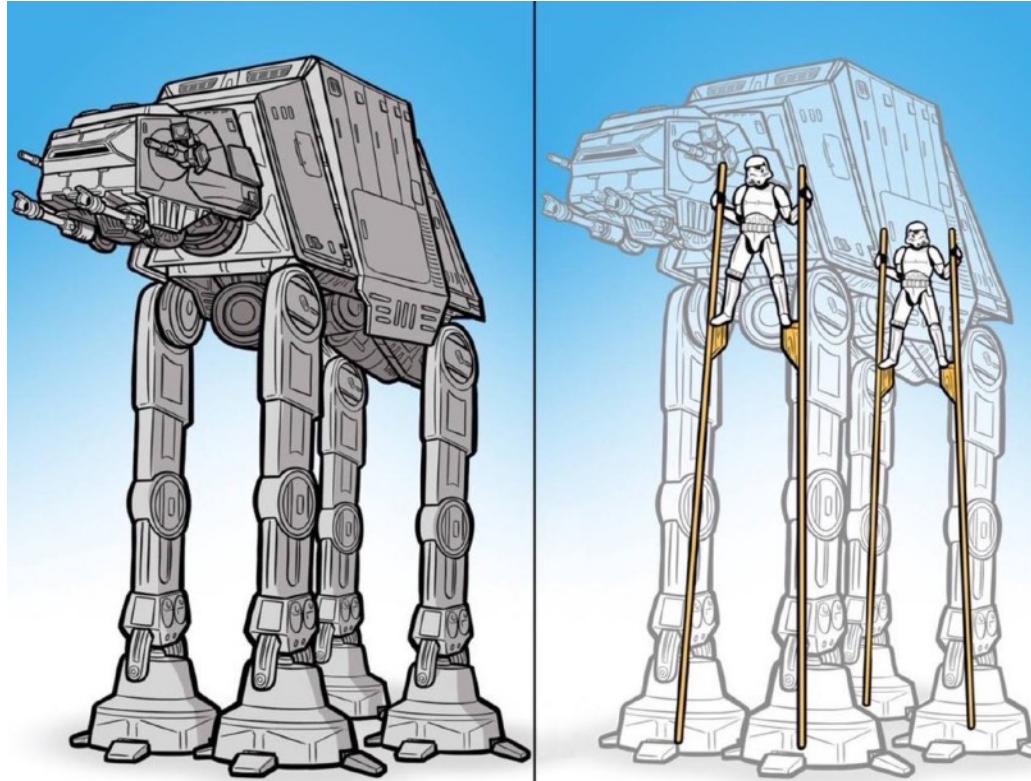
~2000

Today

The Dark Ages

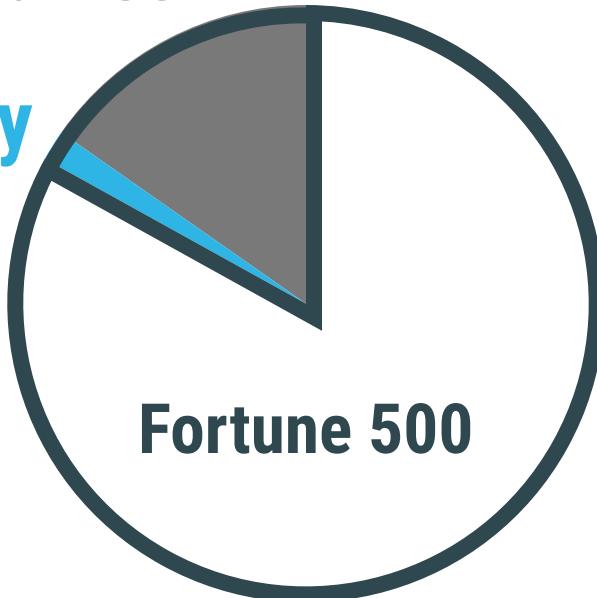




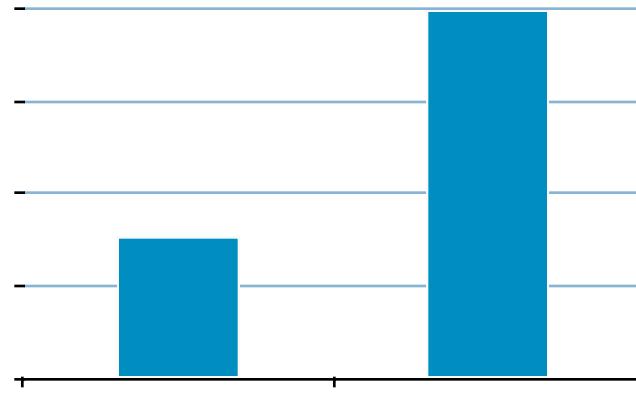


Tech Companies

Good at security



Fortune 500



People
we have

People
we need

Building not breaking

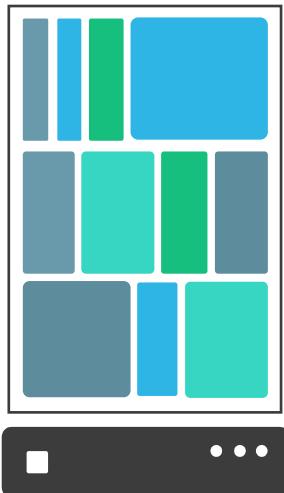
If it doesn't come on by
default, no one will use it



Biased Timeline of the Evolution of Security

~2000

Today



Immutable Infrastructure

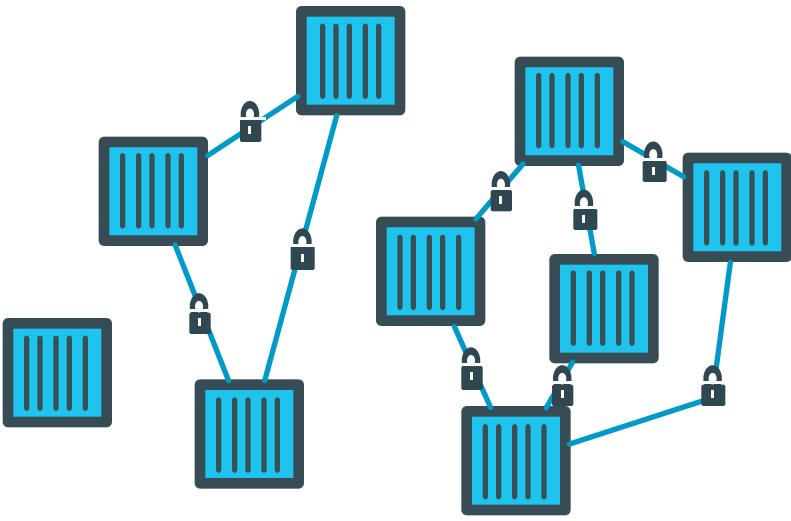
• • •

```
# echo "evil!" > /bin/sh
/bin/sh: can't create /bin/sh: Read-only file system
```

Application Containers



Decentralized Authorization

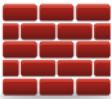


~2000

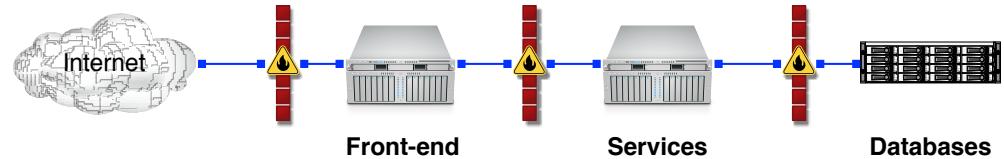
Today



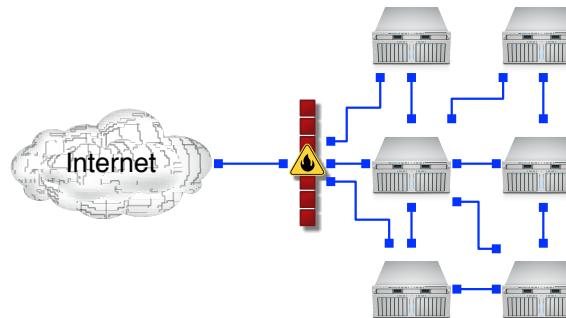
panda



Security
boundary is the
service



VS



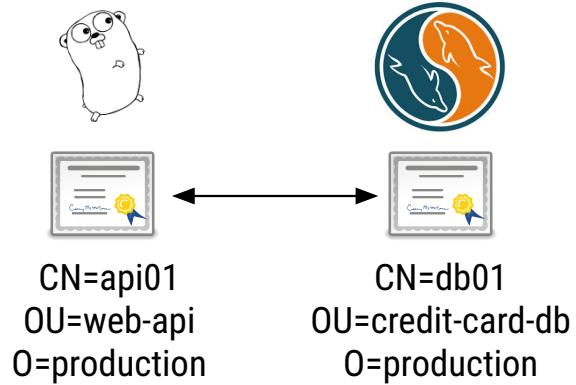
One App Instance, One ID



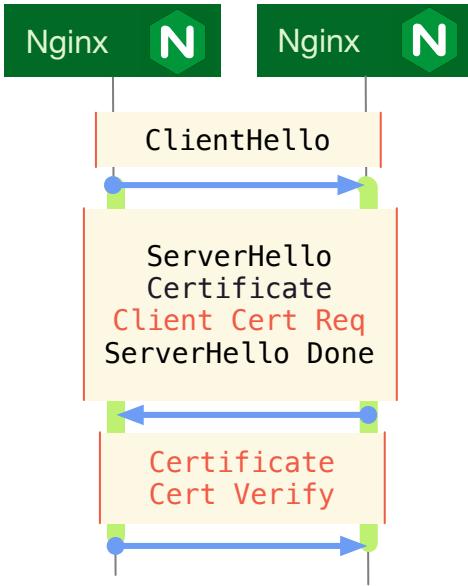
One Node,
One ID



Every service
call is
authorized and
authenticated



mTLS



mTLS - the good

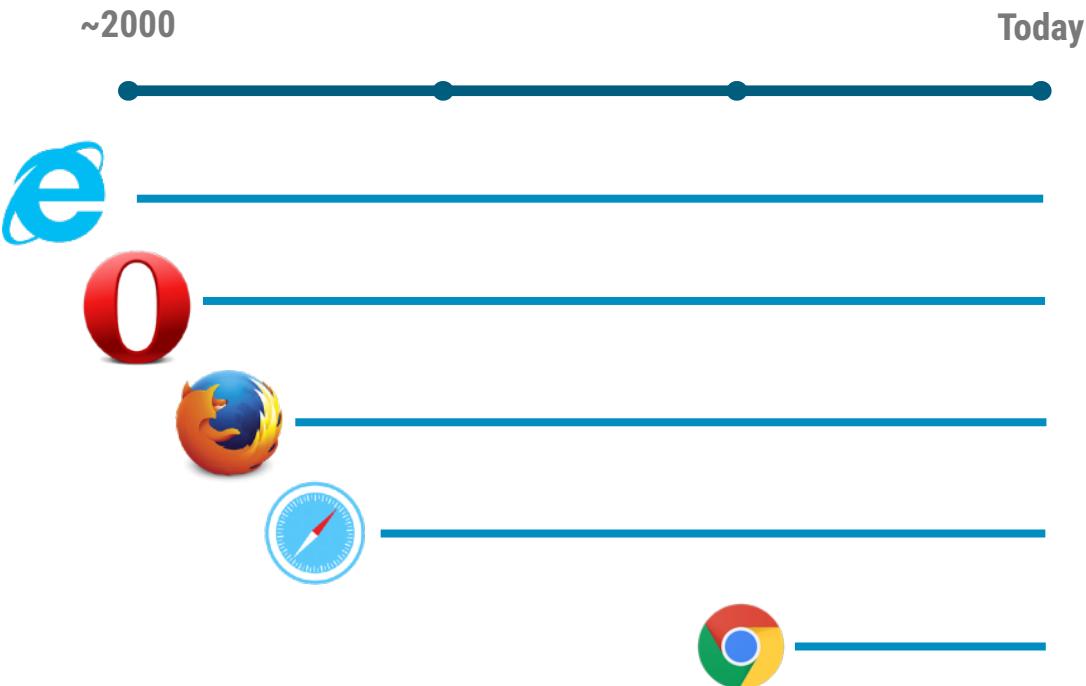
- 1. Key-material stays secret**
- 2. Supported everywhere**
- 3. Authentication and Encryption**

mTLS - the bad

- 1. A LOT of certs**
- 2. Confusing for developers**
- 3. No good revocation story**
- 4. Running a PKI is hard**
- 5. Unforgiving**

Least-privilege resource access

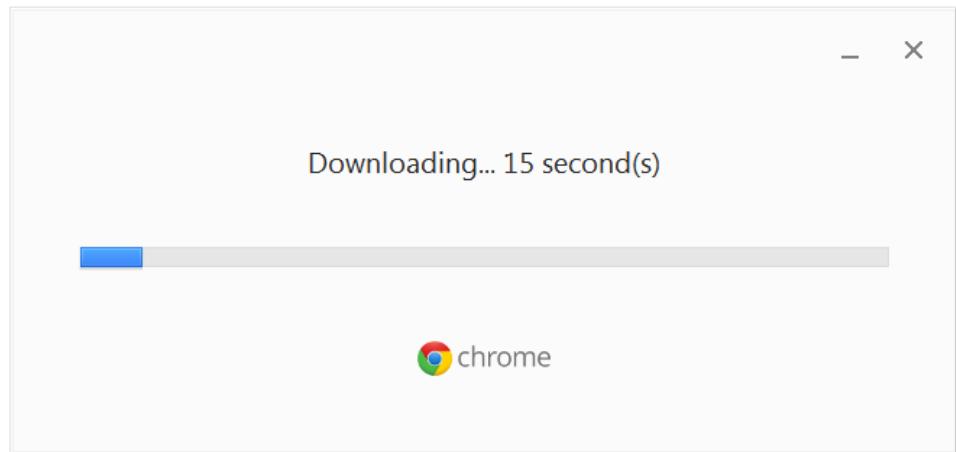
```
[ { "permission":  
    { "method": "GET", "resource": "/user" },  
    "allow": [ "web", "fulfillment", "payments" ] },  
  
  { "permission":  
    { "method": "POST", "resource": "/user" },  
    "allow": [ "signup", "web" ] },  
  
  { "permission":  
    { "method": "DELETE", "resource": "/user/.*" },  
    "allow": [ "web" ]  
  } ]
```



Sandboxing



Transparent Updates



~2000

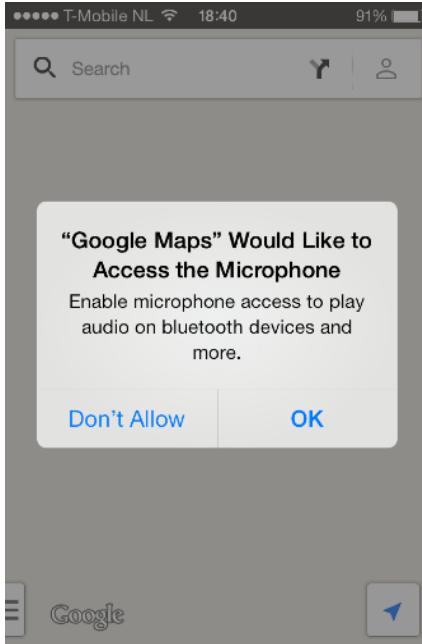
Today

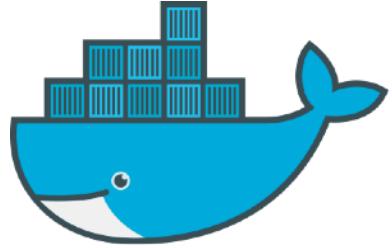


Application Whitelist

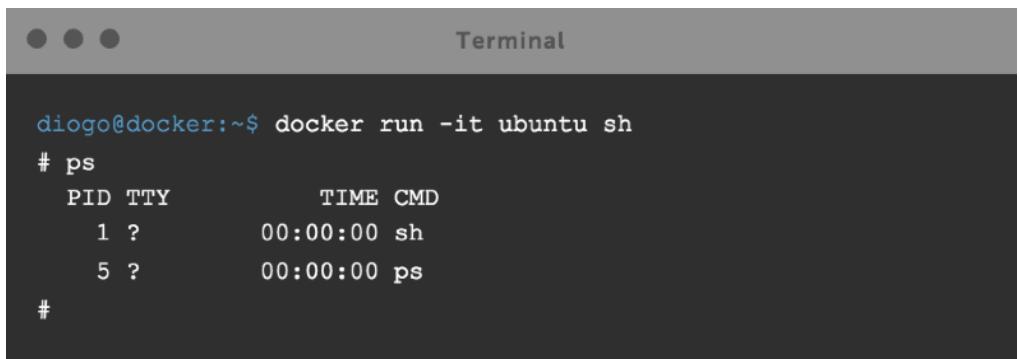


Usable Security





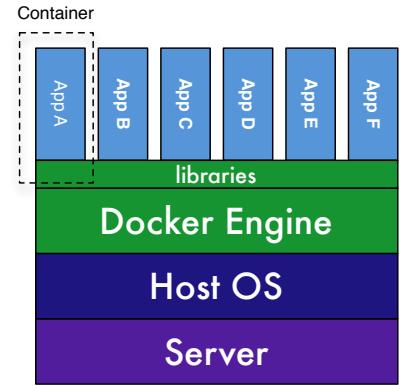
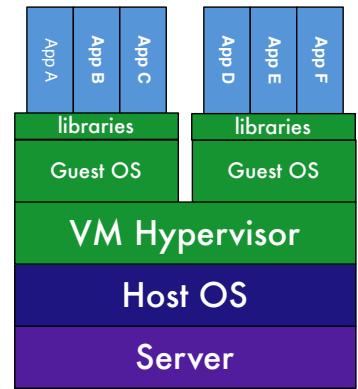
A container is isolated



A screenshot of a Mac OS X terminal window titled "Terminal". The window shows a command-line session on a Ubuntu system within a Docker container. The user has run the command "ps" to list processes. There are two processes listed: one with PID 1, TTY ?, TIME 00:00:00, and CMD sh; and another with PID 5, TTY ?, TIME 00:00:00, and CMD ps. A final "#>" prompt is visible at the bottom.

```
diogo@docker:~$ docker run -it ubuntu sh
# ps
  PID TTY      TIME CMD
    1 ?        00:00:00 sh
    5 ?        00:00:00 ps
#
#
```

Containers VS VMs

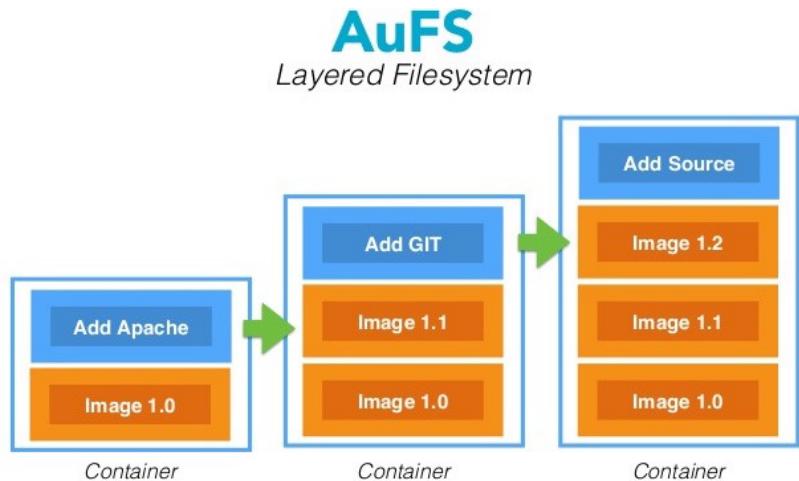


A container is a process

```
root@ubuntunew:/home/diogo# strace -c -f -p 6389
Process 6389 attached
Process 6476 attached
Process 6477 attached
Process 6479 attached
Process 6480 attached
Process 6481 attached
-----
```

% time	seconds	usecs/call	calls	errors	syscall
62.83	0.000747	68	11	6	wait4
7.91	0.000094	6	17		lstat
5.05	0.000060	15	4		getdents64
3.45	0.000041	2	19		open
3.20	0.000038	0	106		read
2.61	0.000031	31	1		nanosleep
2.02	0.000024	2	10		mprotect
1.85	0.000022	1	38		rt_sigaction
1.60	0.000019	0	43		ioctl
1.43	0.000017	1	20		close
1.26	0.000015	0	83		writev

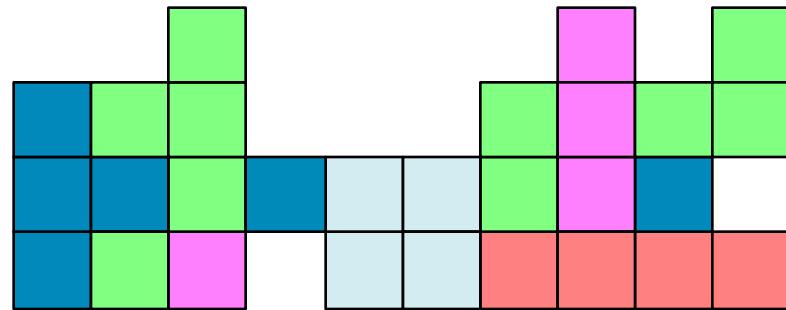
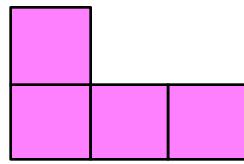
Copy-on-write FS

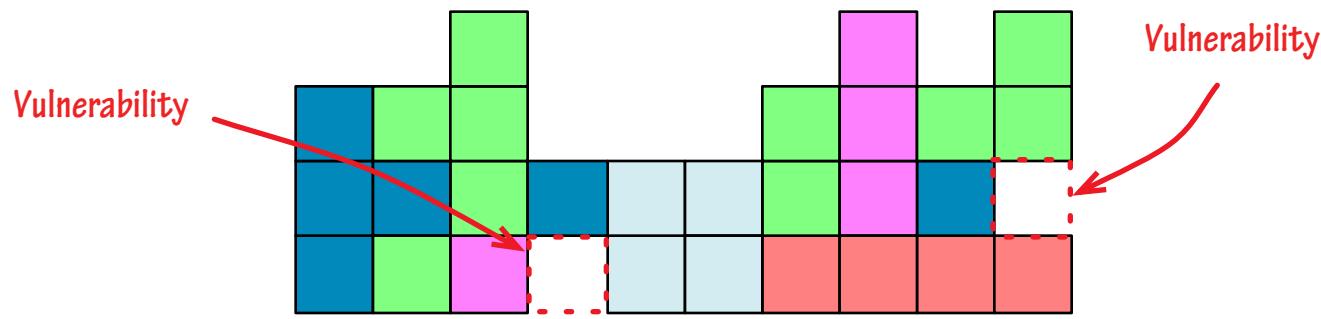
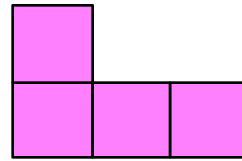


Key Security advantages

- 1. Additional layer of isolation**
- 2. Easier patch management, faster deployments, reduced risk**
- 3. Better visibility into application behavior**

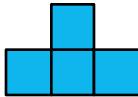
```
● ● ●  
# docker stack deploy --stack-file docker-stack.yml my_app  
Creating service my_app_frontend  
Creating service my_app_backend  
Creating service my_app_db
```



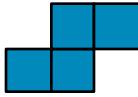


Security Tetrominos

runC



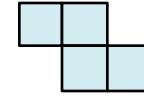
containerD



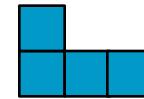
Docker



infraKit



linuxKit



Notary

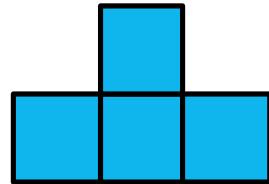


swarmKit



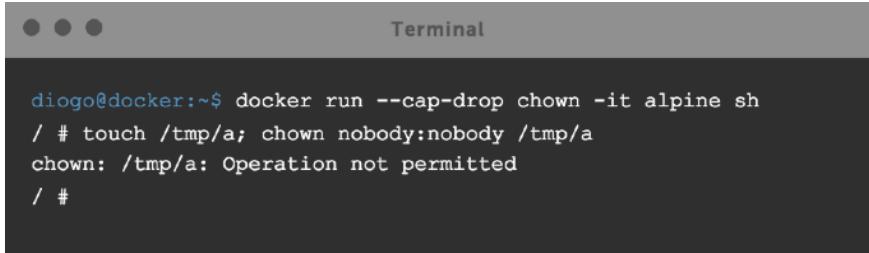
runC

Lightweight universal container runtime



runC

Capabilities



A terminal window titled "Terminal" with three dots at the top left. The window contains the following text:

```
diogo@docker:~$ docker run --cap-drop chown -it alpine sh
/ # touch /tmp/a; chown nobody:nobody /tmp/a
chown: /tmp/a: Operation not permitted
/ #
```

runC

- Namespace Isolation
- Cgroups

Namespaces



Cgroups



runC

Linux Security Modules

1. AppArmor

2. SELinux

3. Smack

4. TOMOYO

5. YAMA

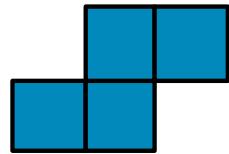
runC

Seccomp-bpf

```
{  
    "defaultAction": "SCMP_ACT_ERRNO",  
    "architectures": [  
        "SCMP_ARCH_X86_64",  
    ],  
    "syscalls": [  
        {  
            "name": "accept",  
            "action": "SCMP_ACT_ALLOW",  
            "args": []  
        },  
        {  
            "name": "accept4",  
            "action": "SCMP_ACT_ALLOW",  
            "args": []  
        },  
        ...  
    ]  
}
```

containerD

Container runtime supervisor



containerD

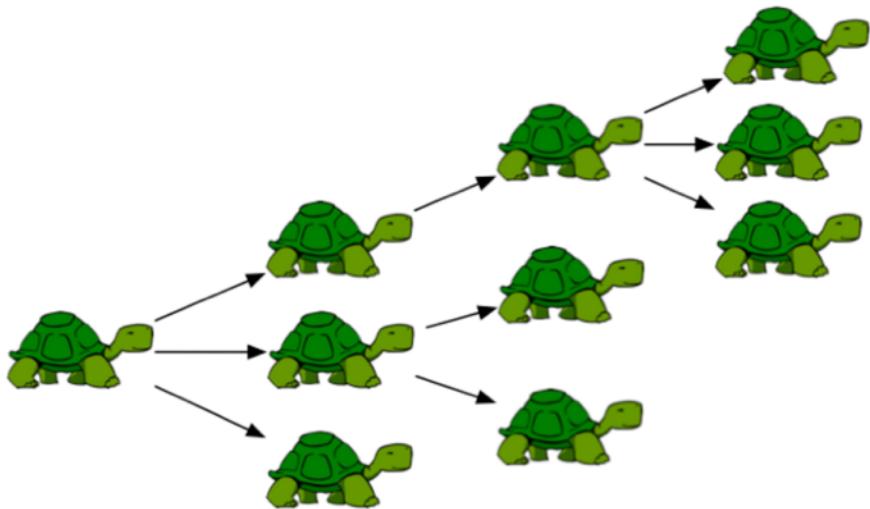
Content
addressable
images

H(blob) ==
3624bbc5e6074c



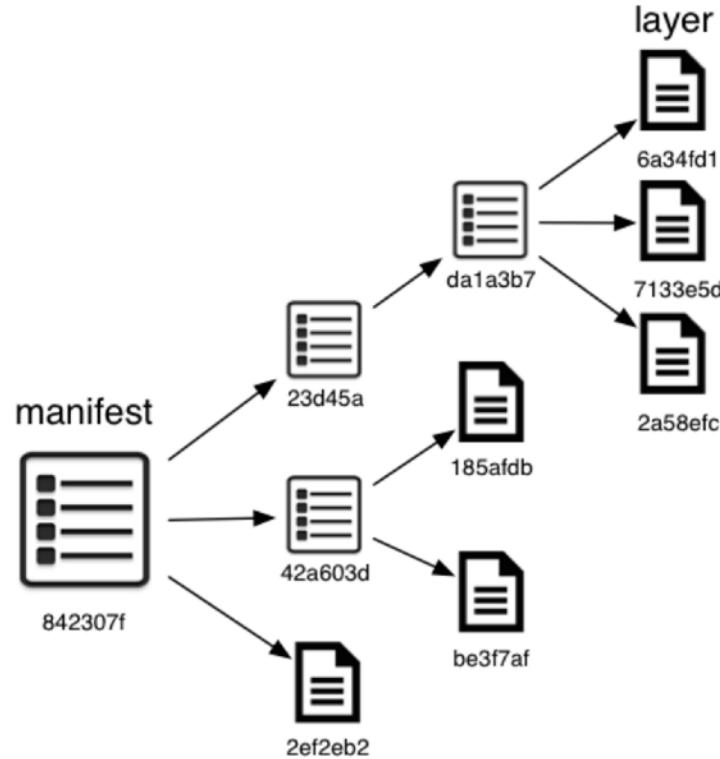
containerD

Turtles all the way
down



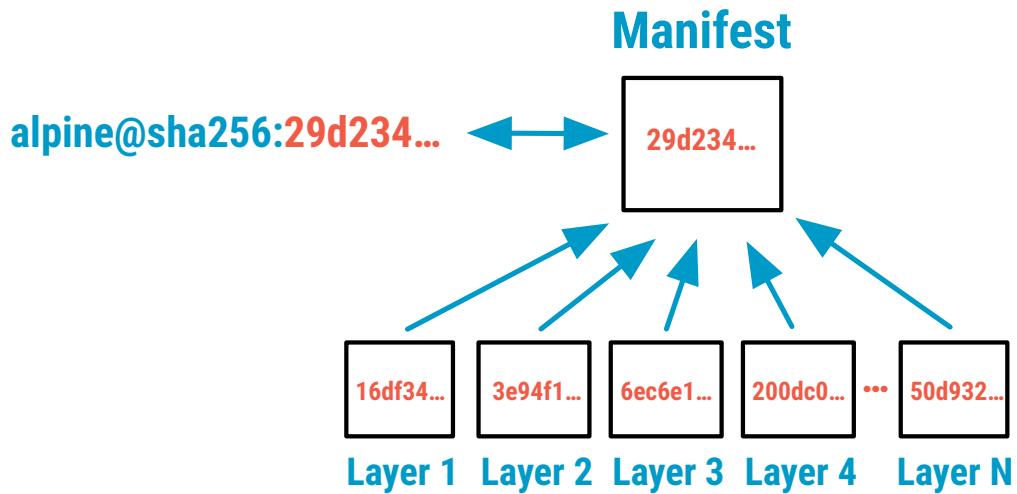
containerD

Images are Merkle
DAGs of layers



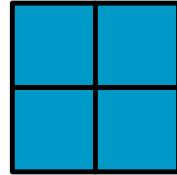
containerD

Content Addressable
Image Pulls



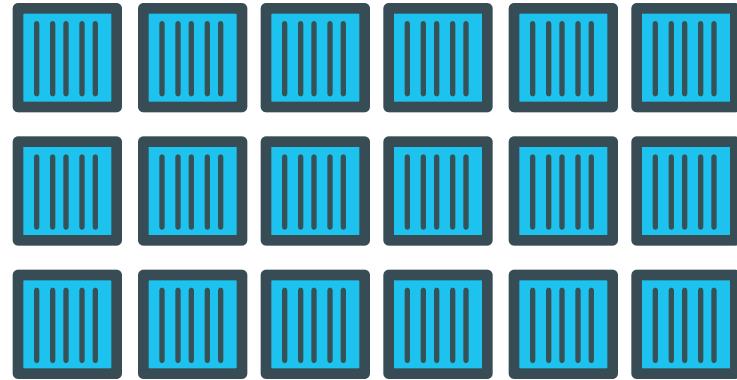
Docker

**Secure-by-default software
container platform**



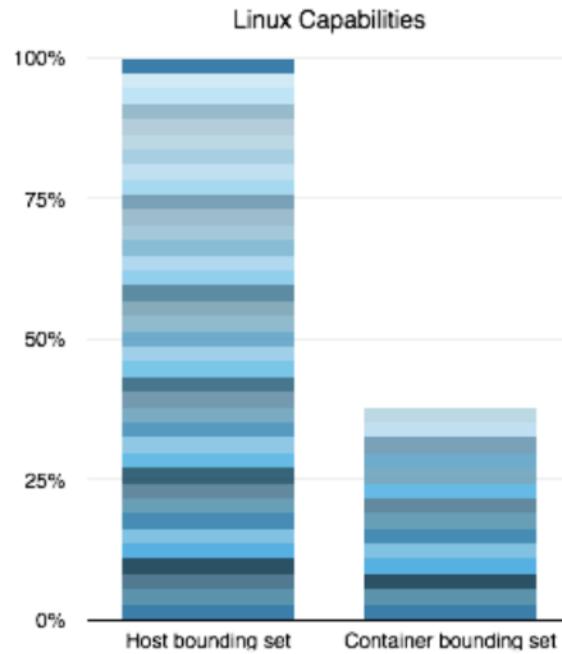
Docker

- SELinux & AppArmor
- Capability Whitelist
- Syscall Whitelist
- Copy-on-write
- No device access



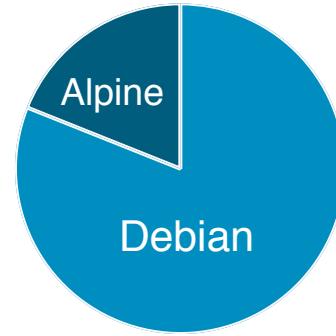
Docker

Less than half the default capabilities



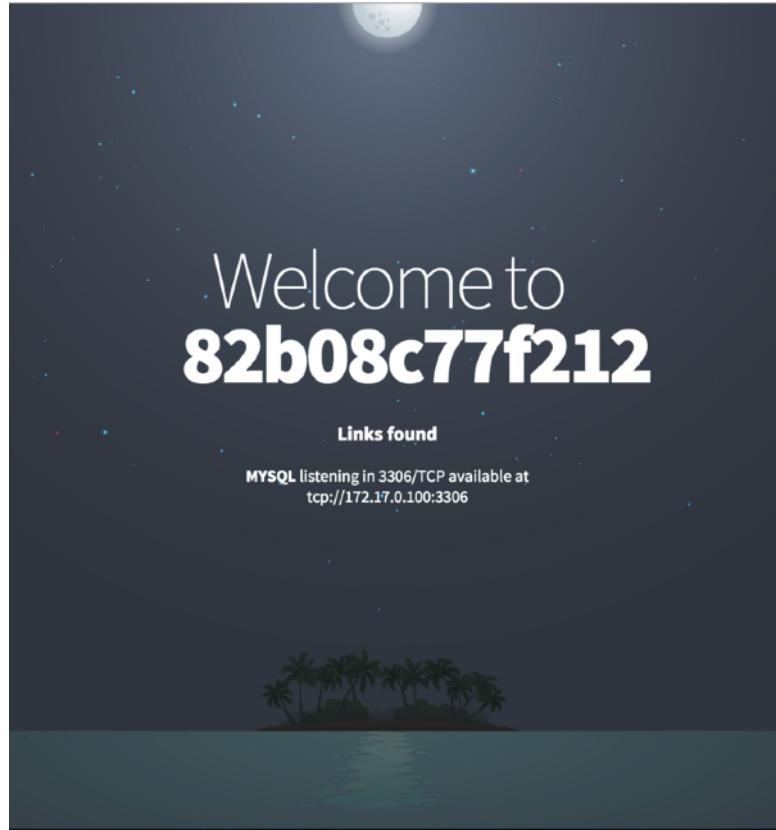
Docker

Minimal Distros



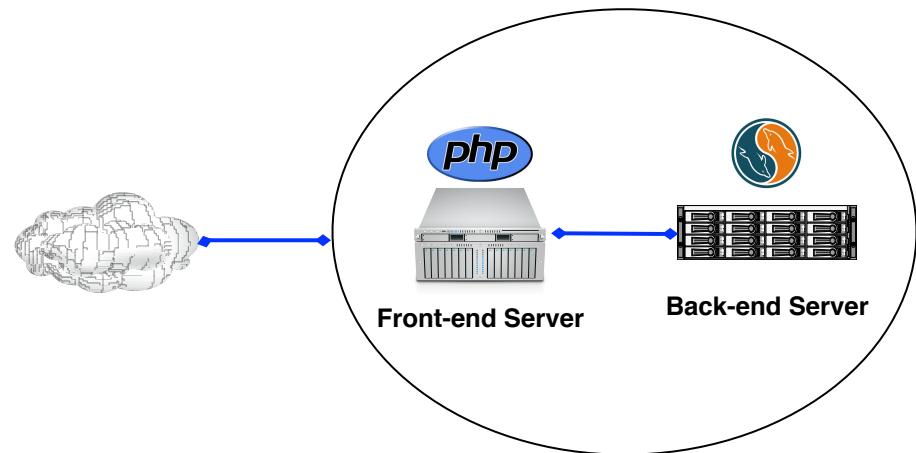
Demo

Our new website



Demo

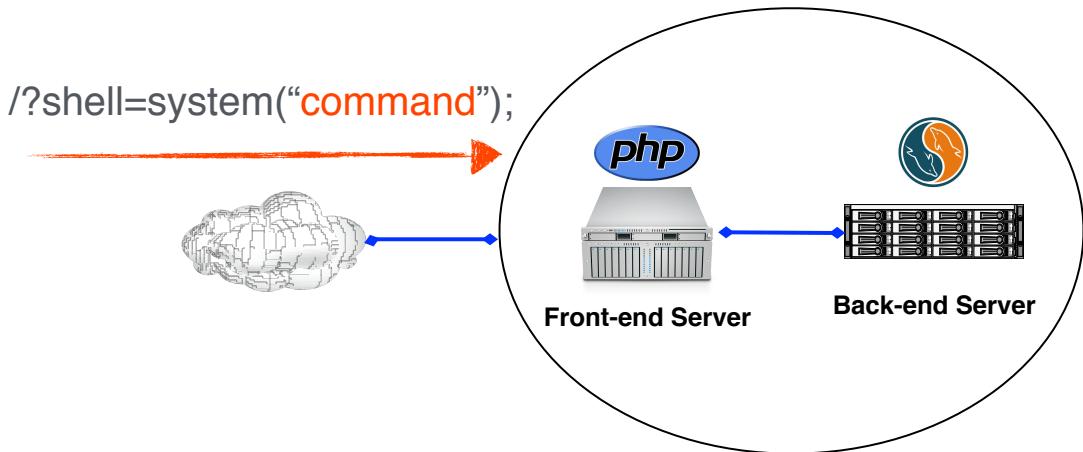
Our new architecture



Demo

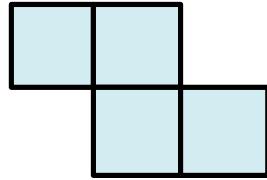
Our new backdoor

```
?shell=system("command");
```



infraKit

Infrastructure independent machine management



infraKit

Platform Agnostic



infraKit

Declarative Updates

• • •

```
# infrakit group commit moby.json
Performing a rolling update on 5 instances
```

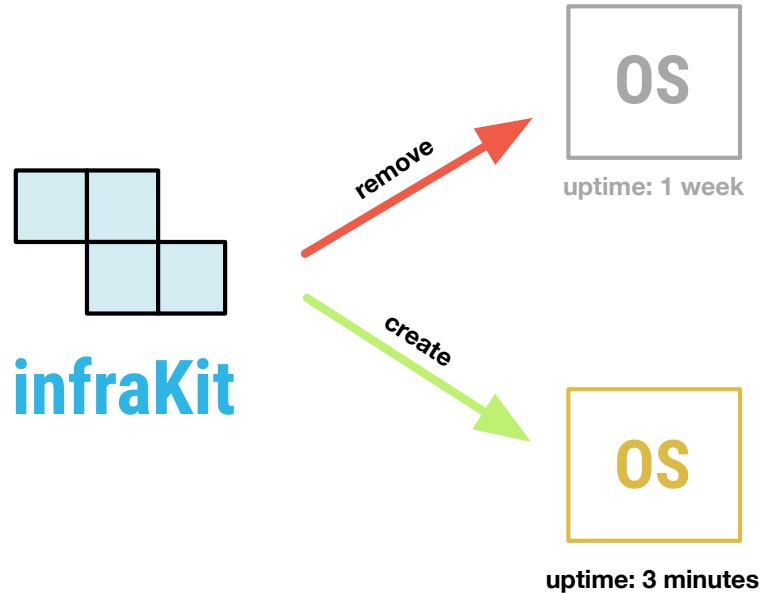
infraKit

Reverse Uptime

10:00

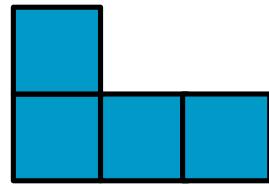
infraKit

Rolling Deploys



linuxKit

A **secure OS builder for your containers**





Immutable Linux OS builder

```
kernel:  
    image: "linuxkit/kernel:4.9.x"  
    cmdline: "console=ttyS0 page_poison=1"  
init:  
    - linuxkit/init:1b8a7e394d...  
onboot:  
    - name: dhcpcd  
        image: "linuxkit/dhcpcd:7d2b8aaaf..."  
        command: ["/sbin/dhcpcd", "--nobackground"]  
trust:  
    org:  
        - linuxkit
```



linuxKit

Minimal Base



linuxKit

Type-safe
System
Daemons



linuxKit

Hardening the kernel

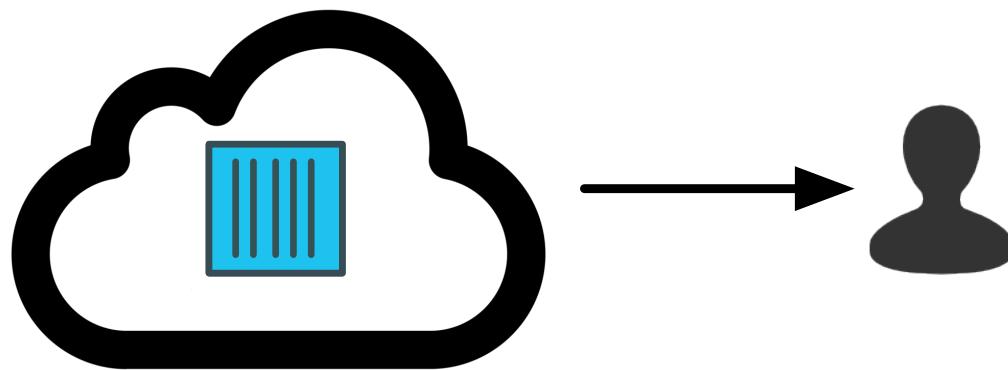


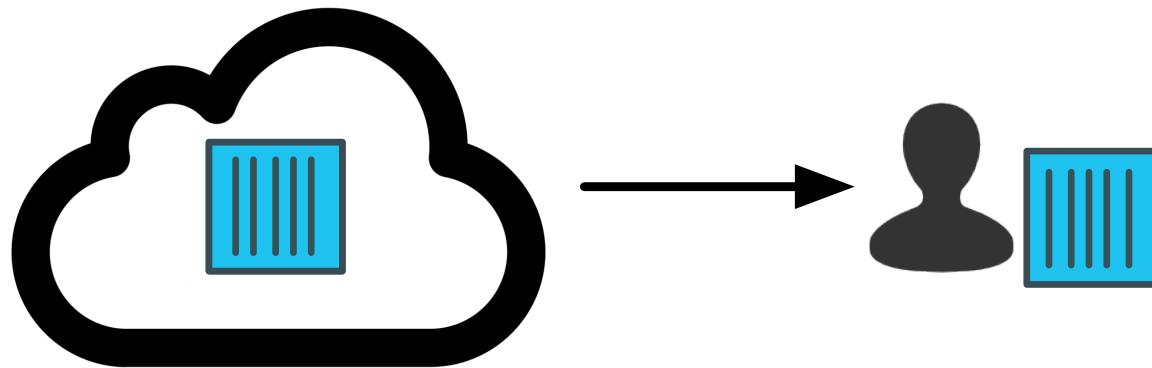
Notary

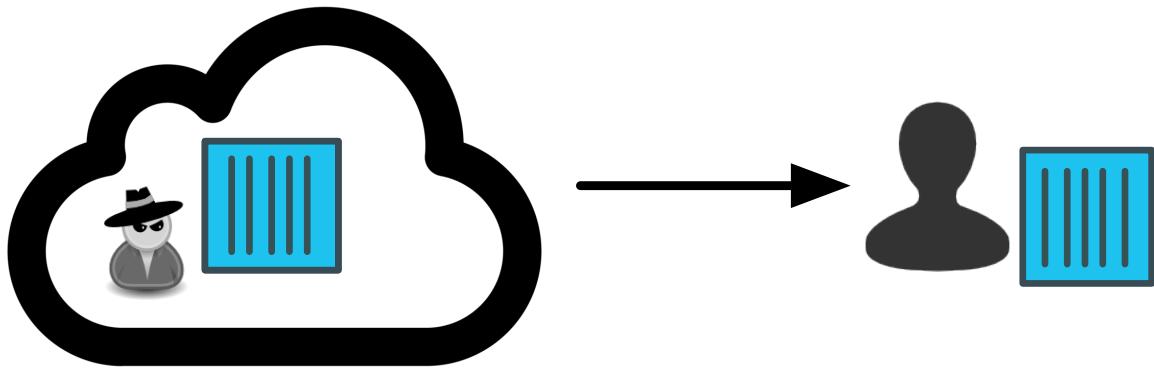
Trusted software delivery



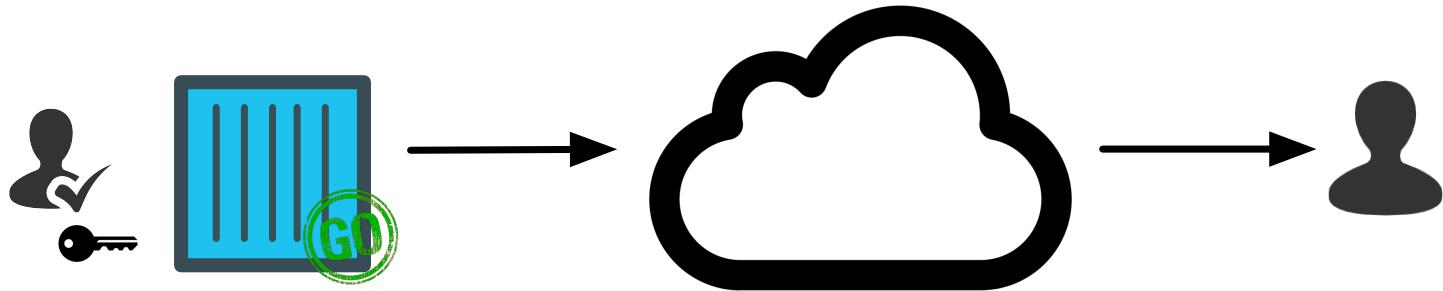
Isn't HTTPS enough?

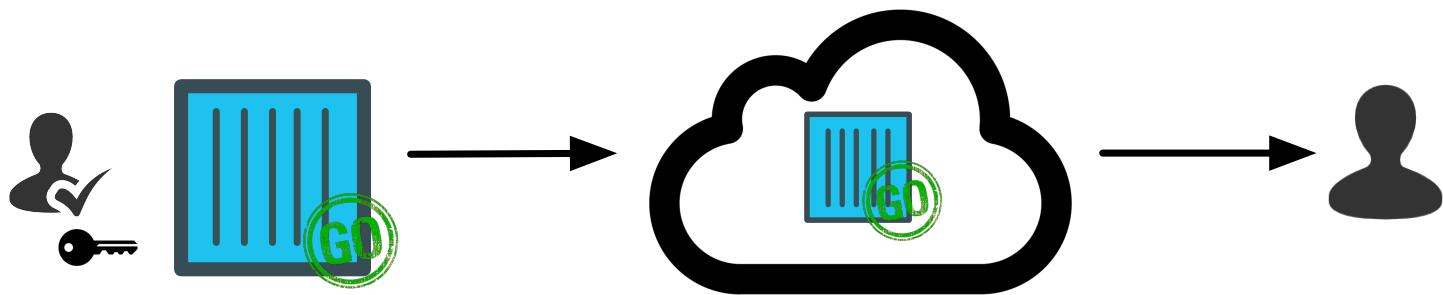


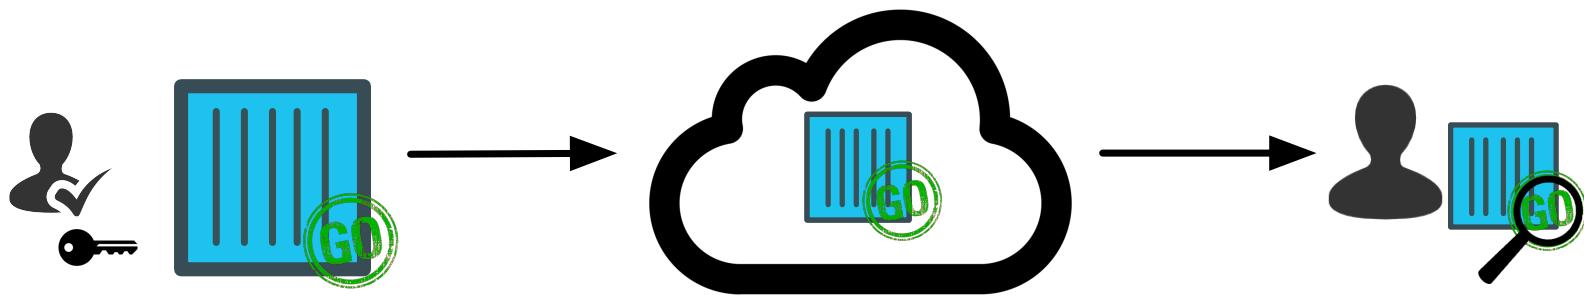


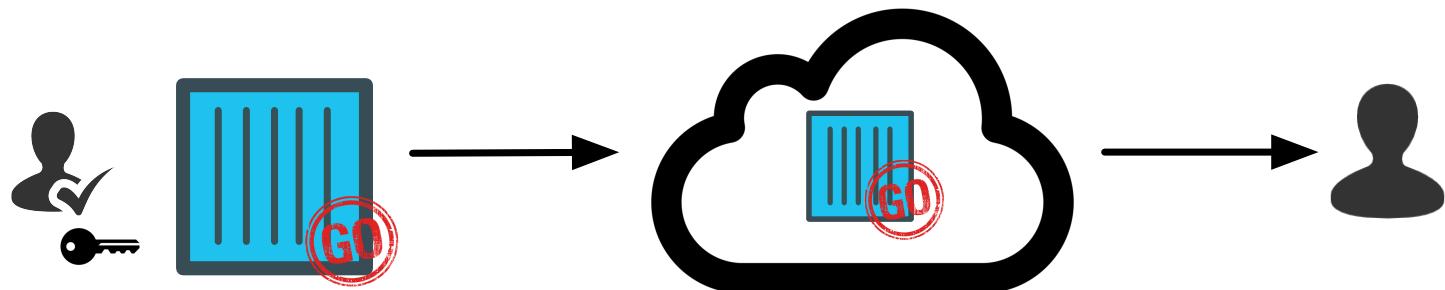


What about GPG?

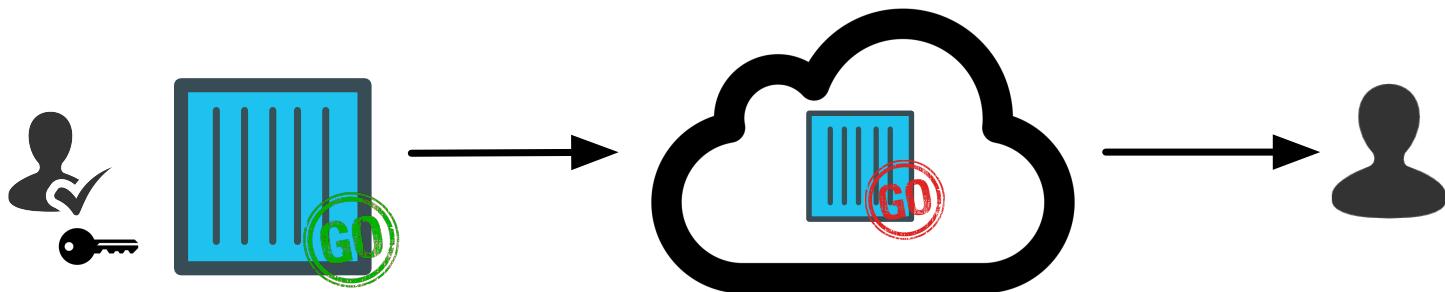


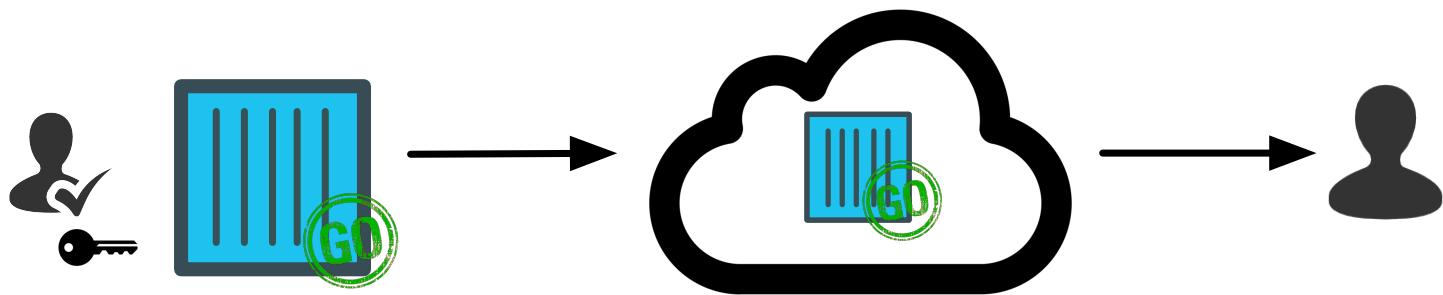


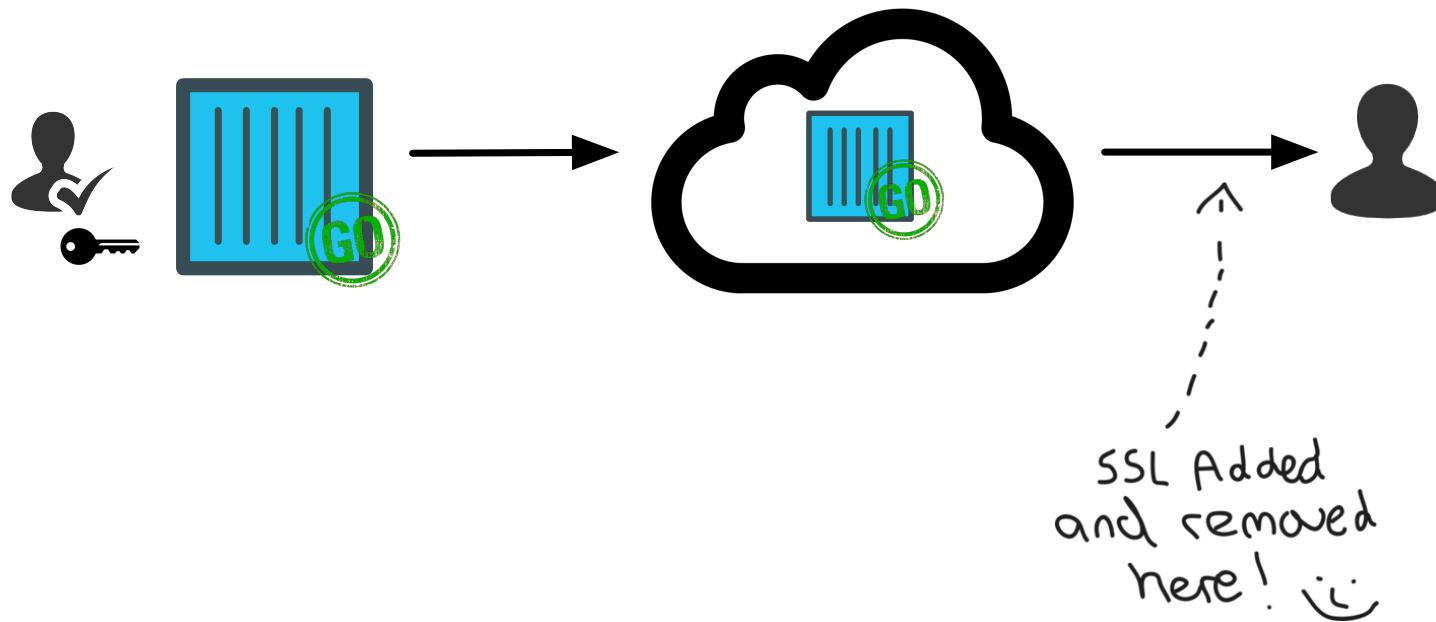


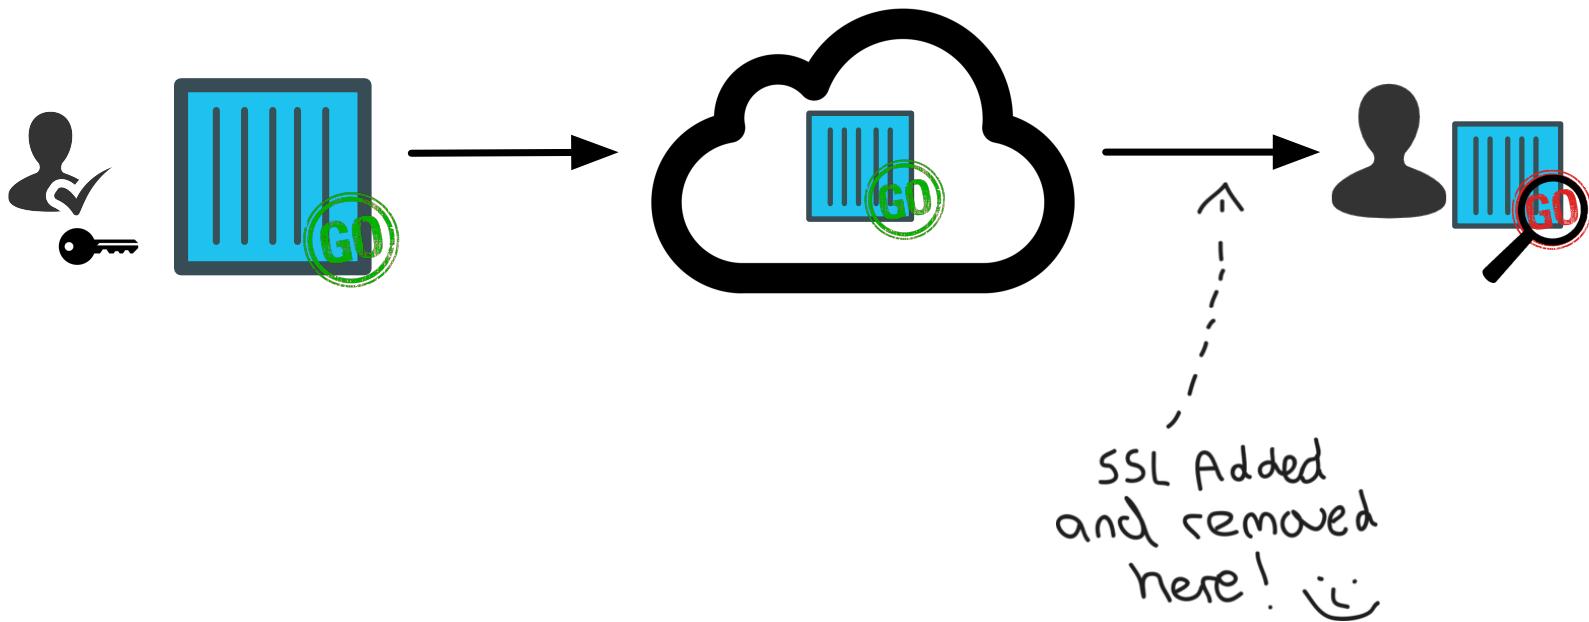


3 months later...









The diagram illustrates the interconnected nature of various software development challenges. At the center is the concept of 'key compromise', which is influenced by and influences several other problematic states:

- Arbitrary Software**: Influences 'key compromise' and is influenced by 'extra dependencies'.
- Mix-and-Match**: Influences 'key compromise' and is influenced by 'extra dependencies'.
- Endless Data**: Influences 'key compromise' and is influenced by 'extra dependencies'.
- Indefinite Freeze**: Influences 'key compromise' and is influenced by 'extra dependencies'.
- Wrong Software**: Influences 'key compromise' and is influenced by 'extra dependencies'.
- Slow Retrieval**: Influences 'key compromise' and is influenced by 'extra dependencies'.
- Malicious Mirrors**: Influences 'key compromise' and is influenced by 'extra dependencies'.
- Rollback**: Influences 'key compromise' and is influenced by 'extra dependencies'.
- Extra Dependencies**: Influences 'key compromise' and is influenced by 'extra dependencies'.

Small text at the bottom right of the diagram area reads: "Source: Software Development Hell".

The diagram illustrates the complexity of secure software distribution by showing various interconnected problems:

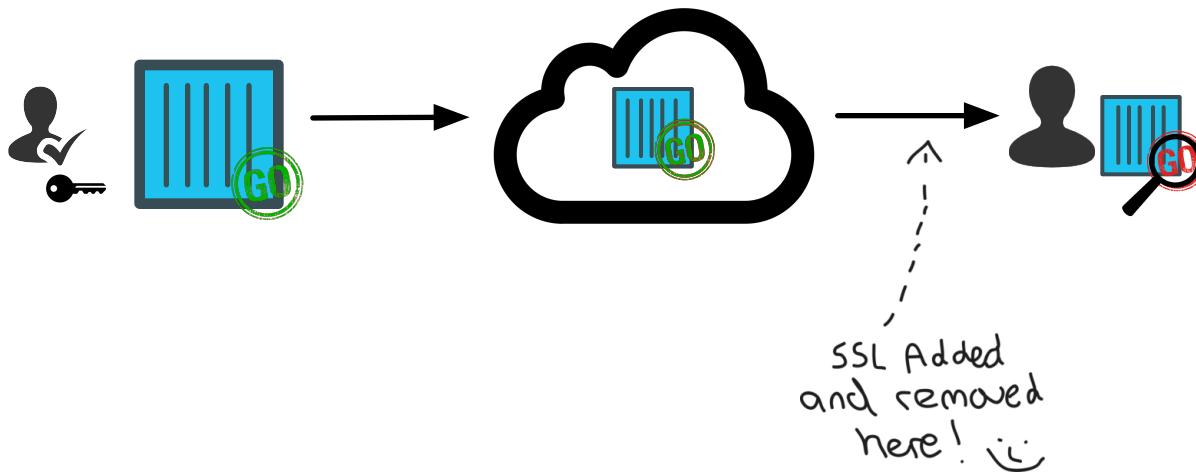
- mix-and-match**: A central issue that overlaps with **arbitrary software**, **endless data**, **key compromise**, **wrong software**, **wrong retrieval**, **indefinite freeze**, **slow retrieval**, **rollback**, and **malicious mirrors**.
- arbitrary software**: Overlaps with **mix-and-match**, **endless data**, **key compromise**, **wrong software**, **wrong retrieval**, **indefinite freeze**, **slow retrieval**, **rollback**, and **malicious mirrors**.
- endless data**: Overlaps with **mix-and-match**, **arbitrary software**, **key compromise**, **wrong software**, **wrong retrieval**, **indefinite freeze**, **slow retrieval**, **rollback**, and **malicious mirrors**.
- key compromise**: Overlaps with **mix-and-match**, **arbitrary software**, **endless data**, **wrong software**, **wrong retrieval**, **indefinite freeze**, **slow retrieval**, **rollback**, and **malicious mirrors**.
- wrong software**: Overlaps with **mix-and-match**, **arbitrary software**, **endless data**, **key compromise**, **wrong retrieval**, **indefinite freeze**, **slow retrieval**, **rollback**, and **malicious mirrors**.
- wrong retrieval**: Overlaps with **mix-and-match**, **arbitrary software**, **endless data**, **key compromise**, **wrong software**, **indefinite freeze**, **slow retrieval**, **rollback**, and **malicious mirrors**.
- indefinite freeze**: Overlaps with **mix-and-match**, **arbitrary software**, **endless data**, **key compromise**, **wrong software**, **wrong retrieval**, **slow retrieval**, **rollback**, and **malicious mirrors**.
- slow retrieval**: Overlaps with **mix-and-match**, **arbitrary software**, **endless data**, **key compromise**, **wrong software**, **wrong retrieval**, **indefinite freeze**, **rollback**, and **malicious mirrors**.
- rollback**: Overlaps with **mix-and-match**, **arbitrary software**, **endless data**, **key compromise**, **wrong software**, **wrong retrieval**, **indefinite freeze**, **slow retrieval**, and **malicious mirrors**.
- malicious mirrors**: Overlaps with **mix-and-match**, **arbitrary software**, **endless data**, **key compromise**, **wrong software**, **wrong retrieval**, **indefinite freeze**, **slow retrieval**, and **rollback**.

“A software update system is secure if it can be sure that it knows about the latest available updates in a timely manner, any files it downloads are the correct files, and no harm results from checking or downloading files.”

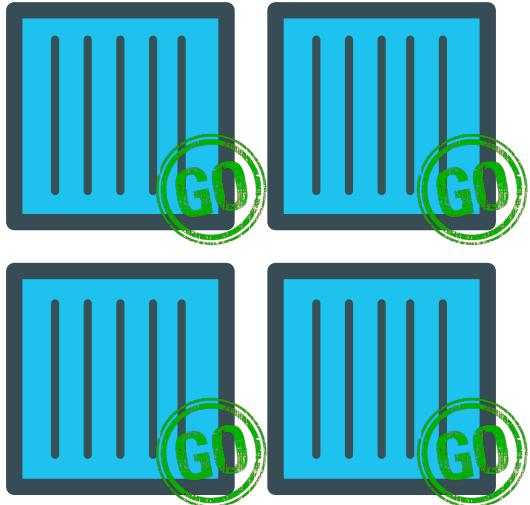
- The Update Framework

- 1. Freshness**
- 2. Signed Collections**
- 3. Key Hierarchy**
- 4. Transparent Key Rotation**
- 5. Threshold Signing**

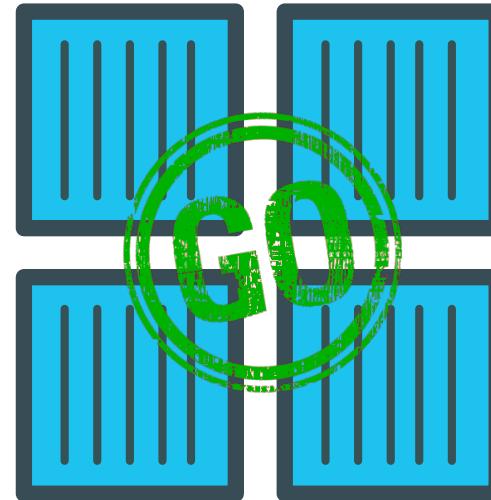
Freshness



Signed Collections

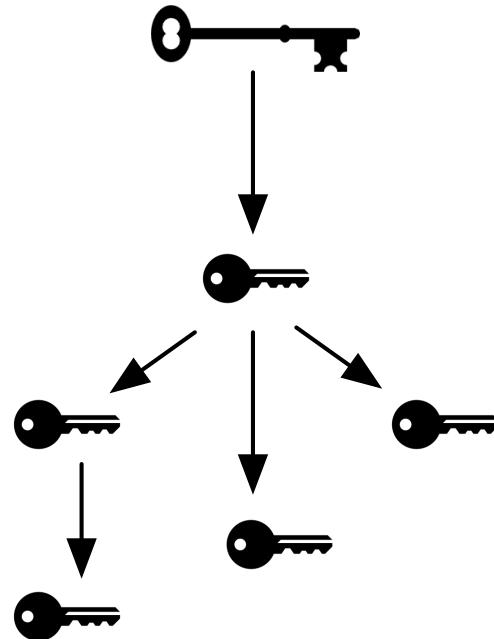


Signed Packages

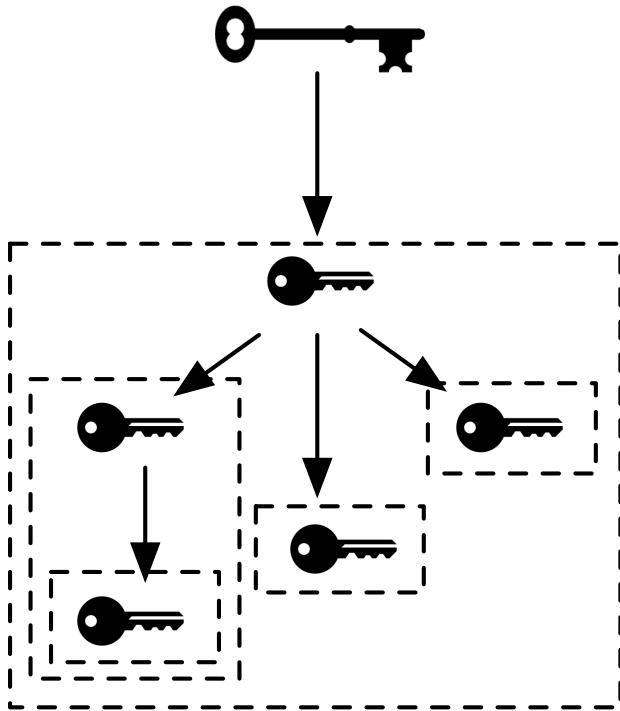


Signed Collection

Key Hierarchy



Key Hierarchy

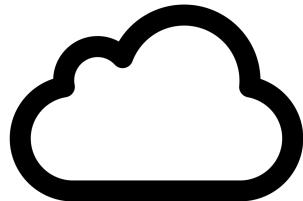


Key Hierarchy

Short
Expiry



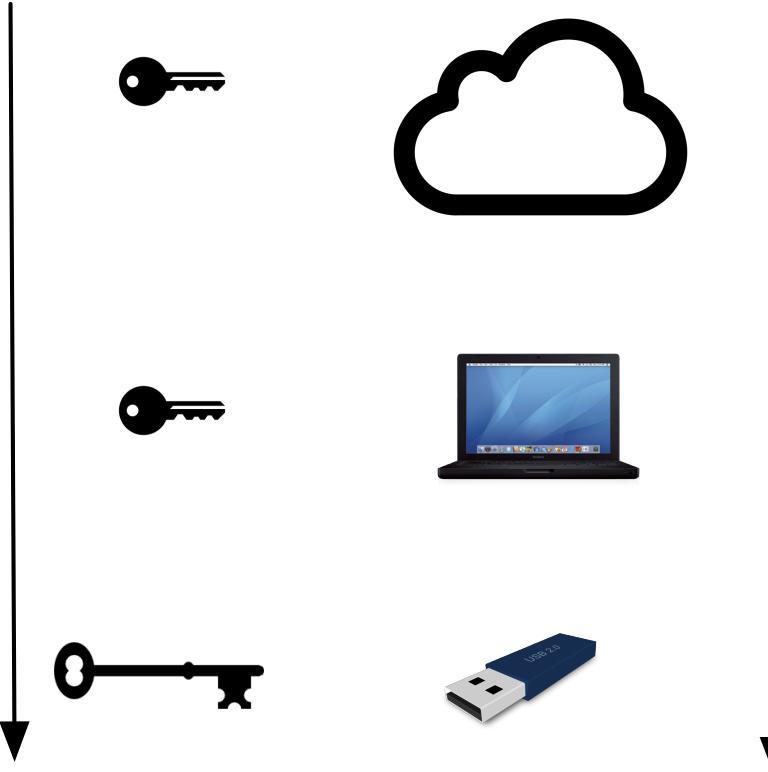
Less
Sensitive



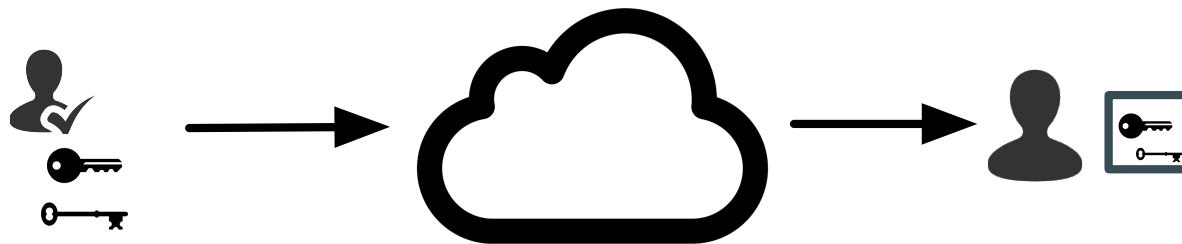
Long
Expiry



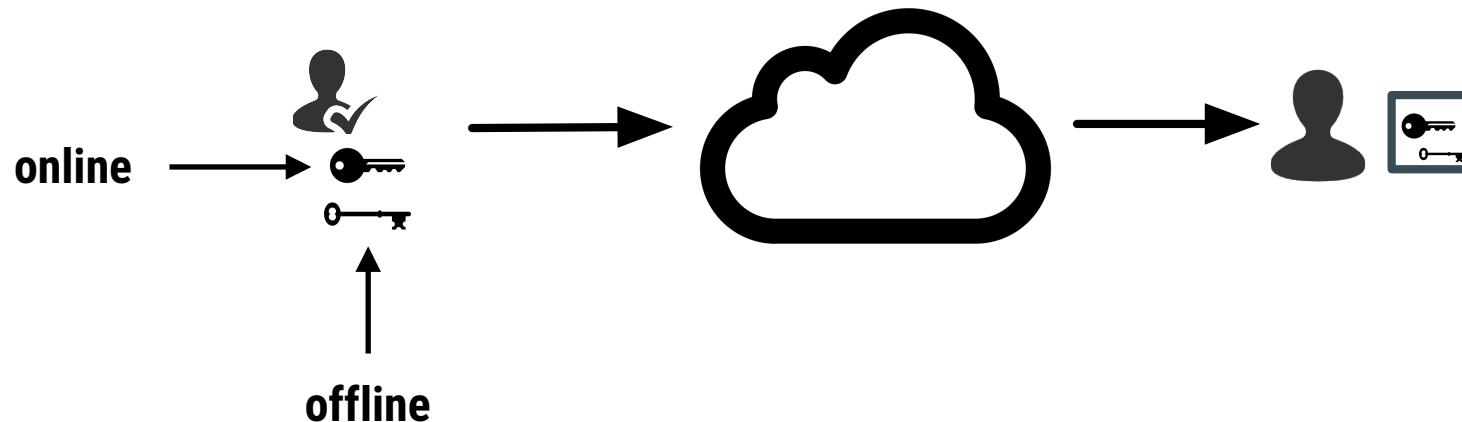
More
Sensitive



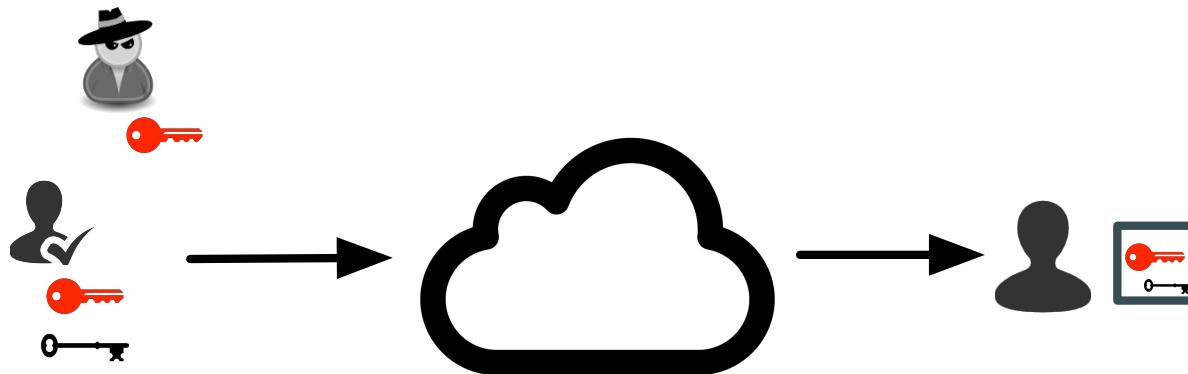
Transparent Key Rotation



Transparent Key Rotation



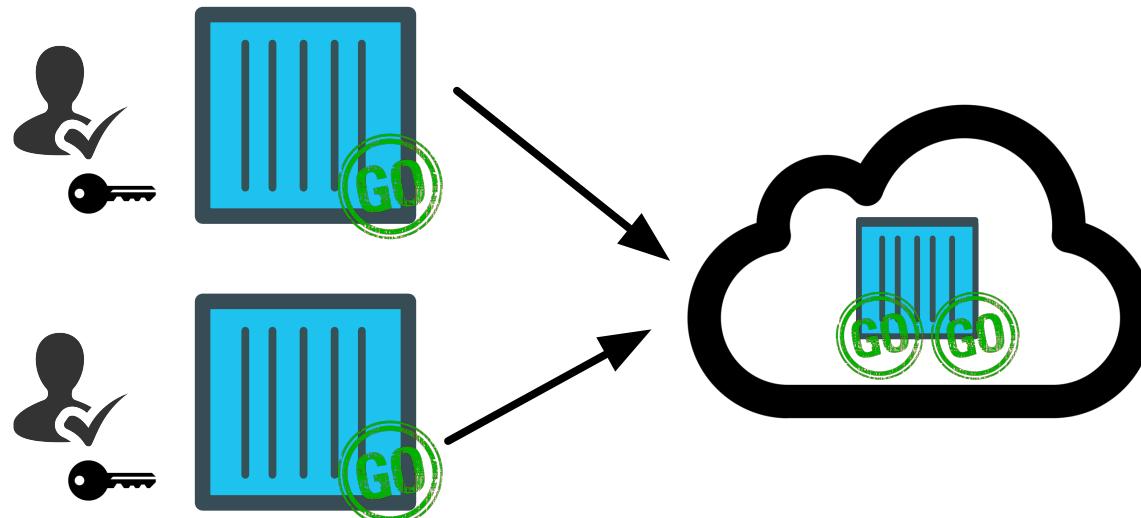
Transparent Key Rotation



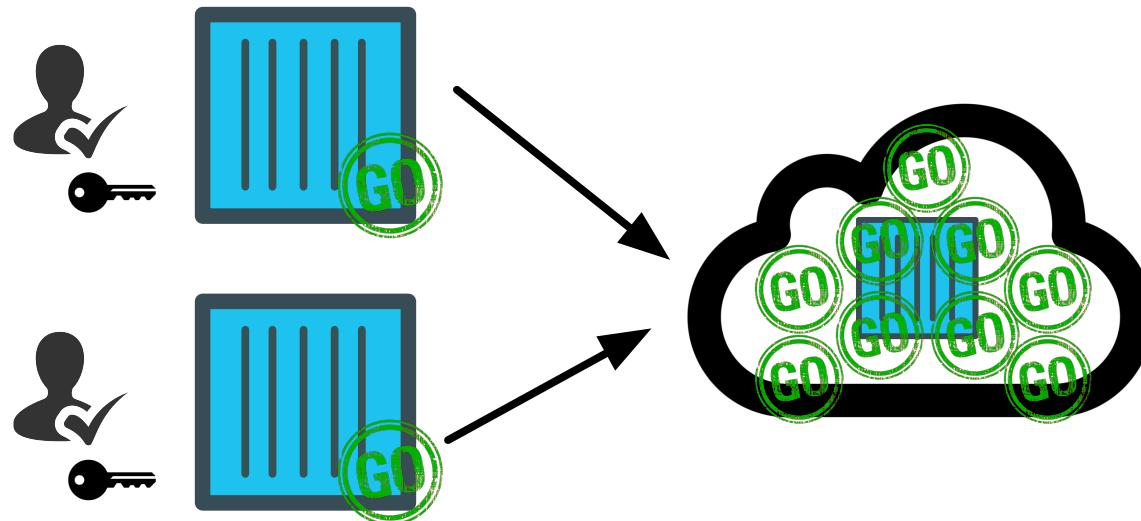
Transparent Key Rotation



Threshold Signing



Threshold Signing



The Update Framework (**TUF**)

theupdateframework.com

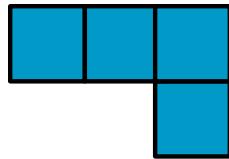
Notary

Cryptographic
Name Resolution



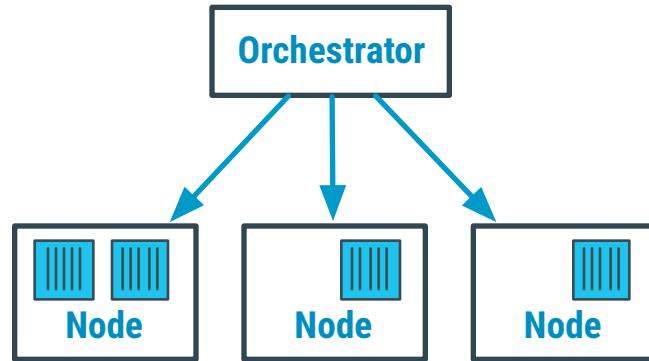
swarmKit

Least-privilege container orchestrator



swarmKit

Multiple
nodes



swarmKit

Secure Node Introduction



swarmKit

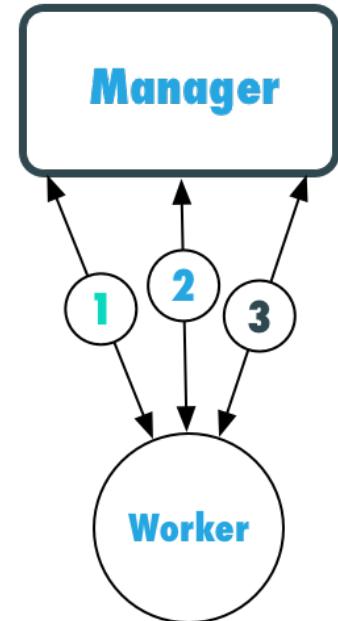
Cryptographic Node Identity

```
$ openssl x509 -in  
/var/lib/docker/swarm/certificates/swarm-node.crt  
-text  
Cert  
...  
Swarm ID  
Issuer: CN=swarm-ca  
Validity  
Not Before: Mar  9 15:21:00 2017 GMT  
Not After : Jun  7 16:21:00 2017 GMT  
Subject: O=lgz5xj1eqg... OU=swarm-manager, CN=o...  
Node Role  
...  
X509v3 Subject Alternative Name:  
DNS:swarm-manager, DNS:o... , DNS:swarm-ca  
...  
-----BEGIN CERTIFICATE-----  
MIICNDCCAdugAwIBAgIUCoRaj23j4h5  
...
```

swarmKit

Bootstrap

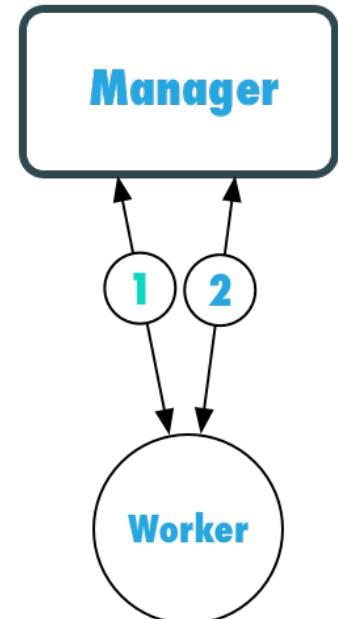
1. Retrieve and validate Root CA Public key material.
2. Submit new CSR along with secret token.
3. Retrieve the signed certificate.



swarmKit

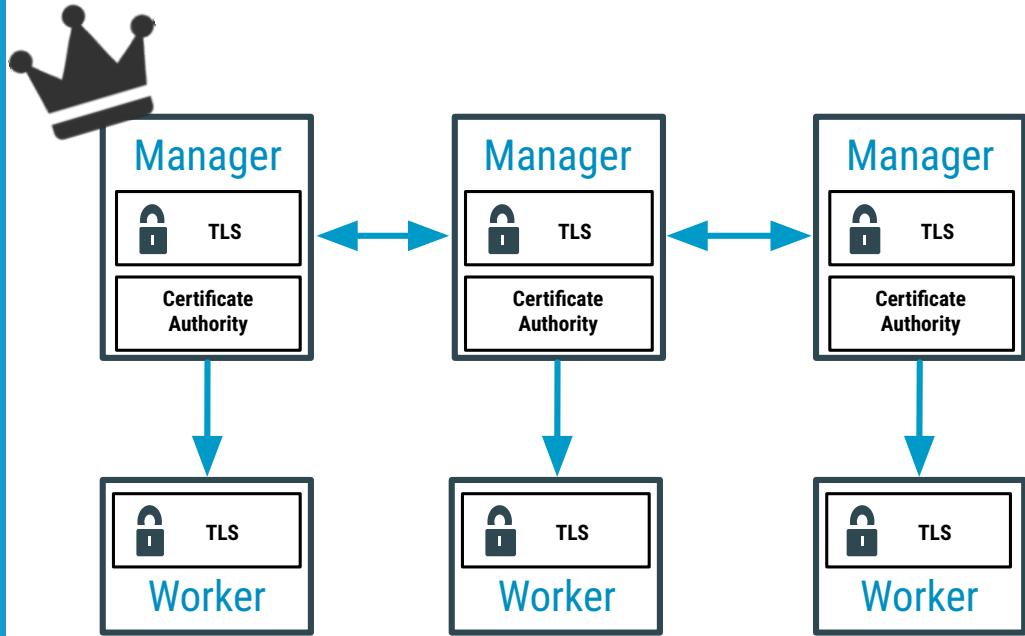
Automatic Certificate Rotation

1. Submit new CSR using old key-pair.
2. Retrieve the new signed certificate.



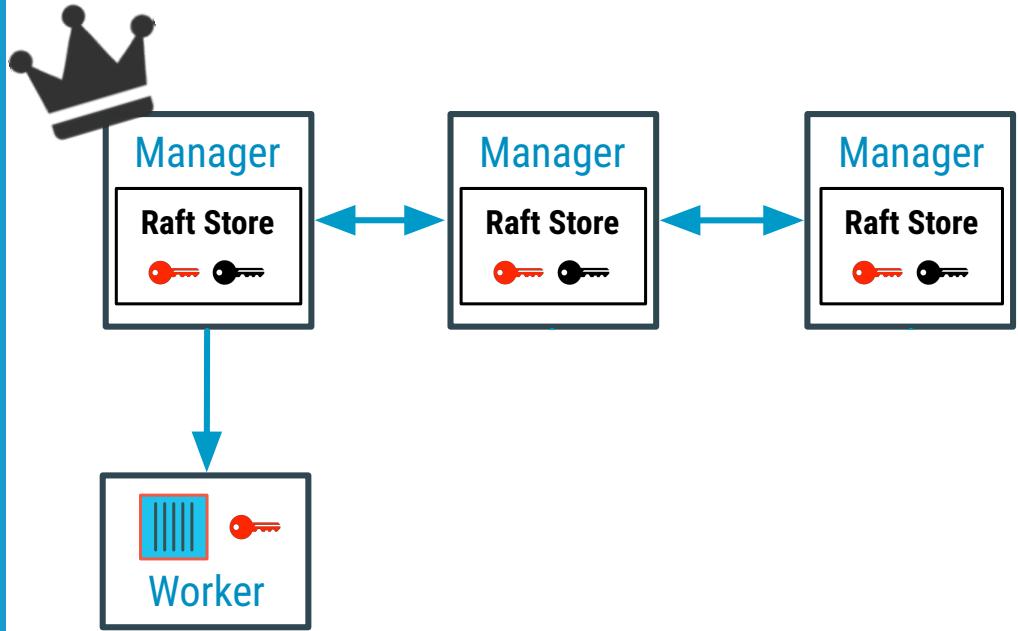
swarmKit

MTLS Between
All Nodes



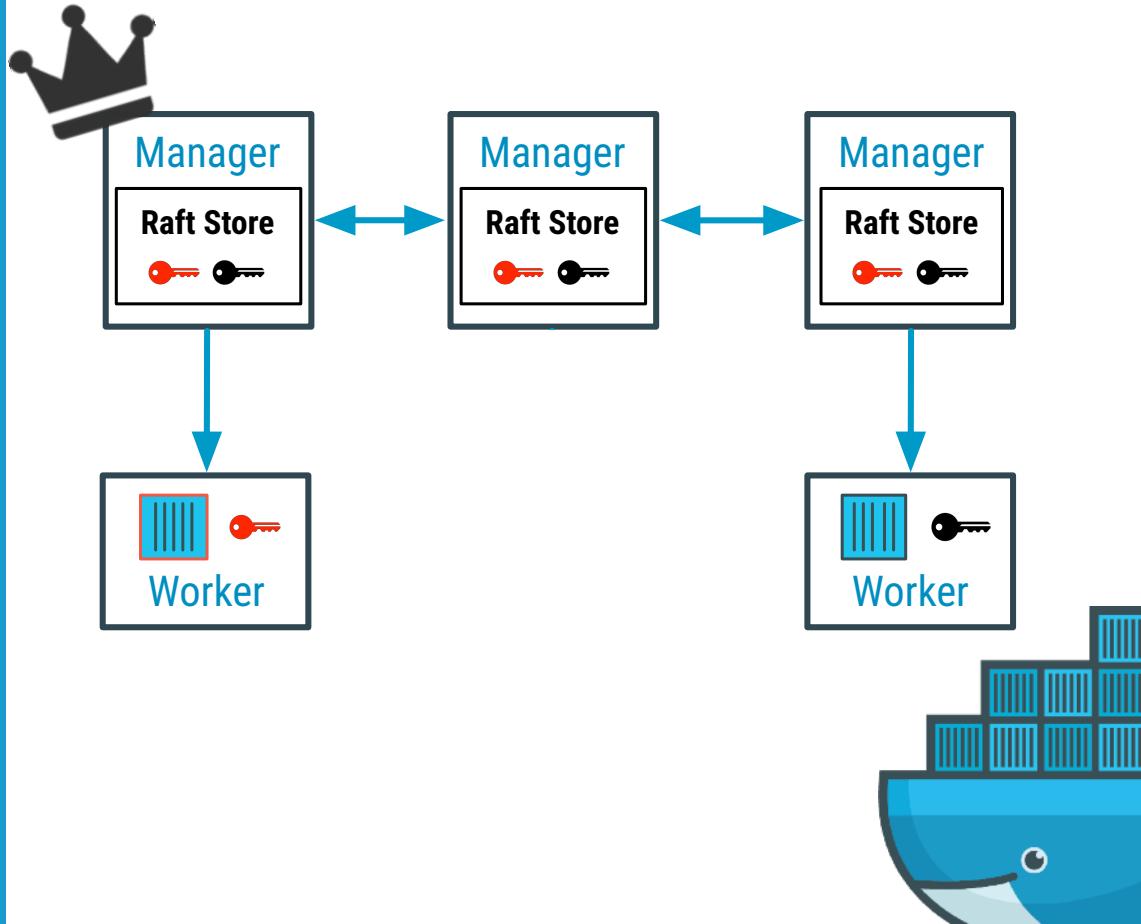
swarmKit

Least-privilege
Secret
Distribution



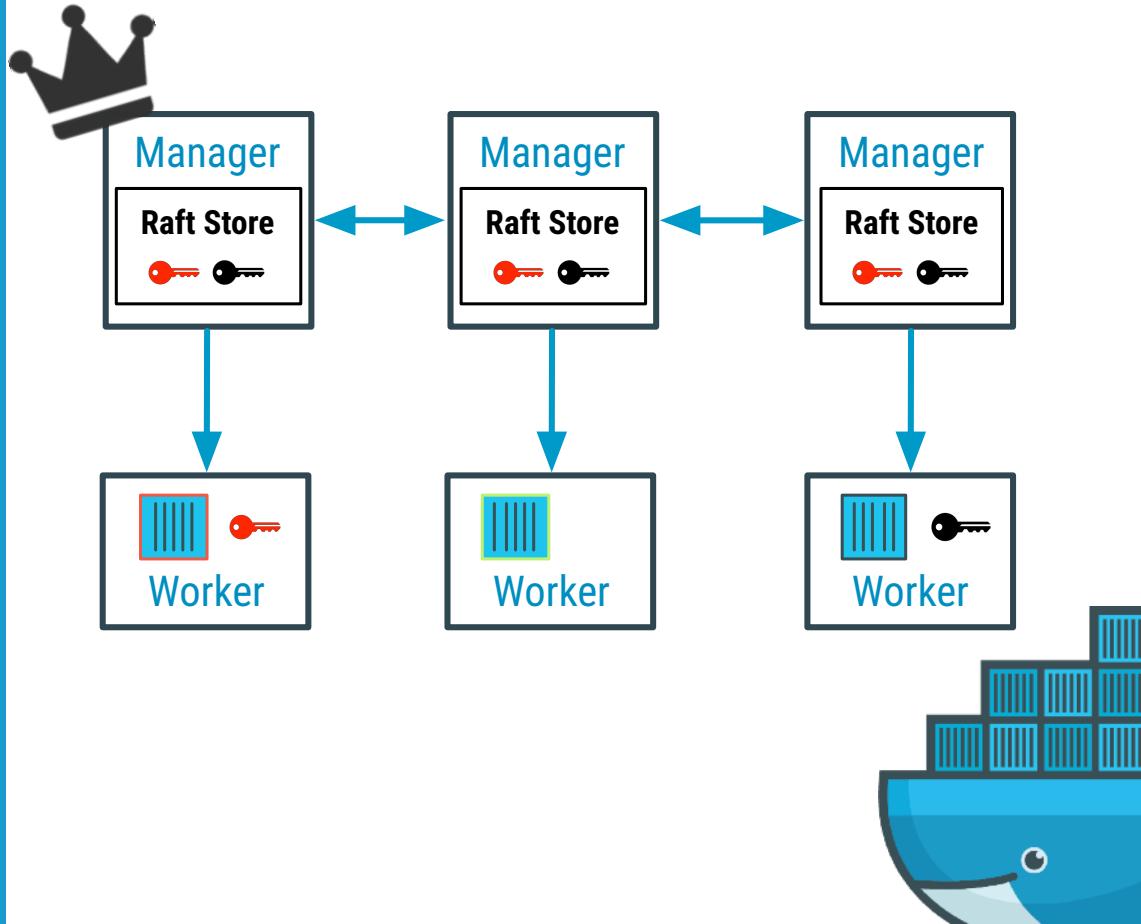
swarmKit

Least-privilege
Secret
Distribution



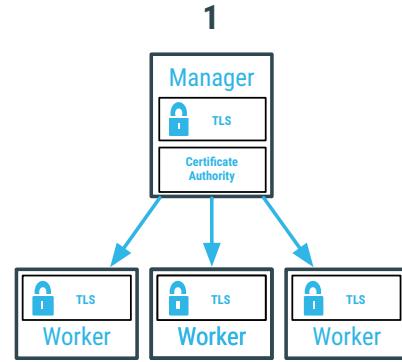
swarmKit

Least-privilege
Secret
Distribution



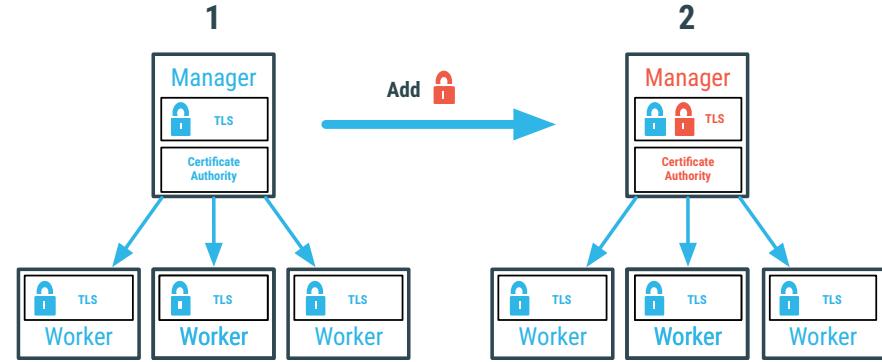
swarmKit

Transparent Root Rotation



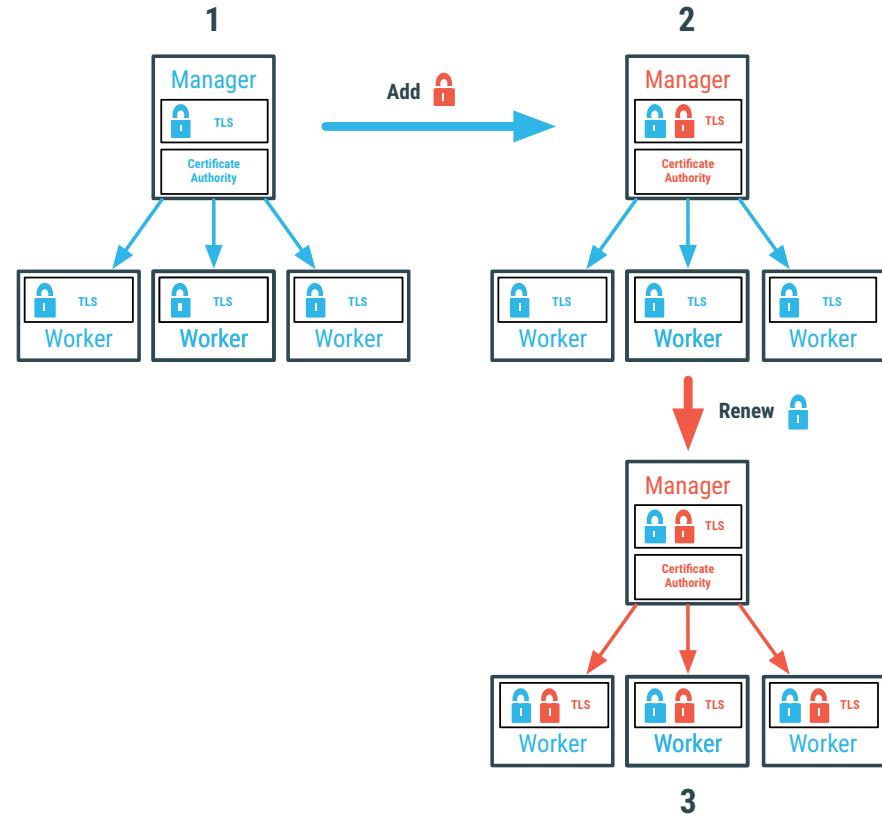
swarmKit

Transparent Root Rotation



swarmKit

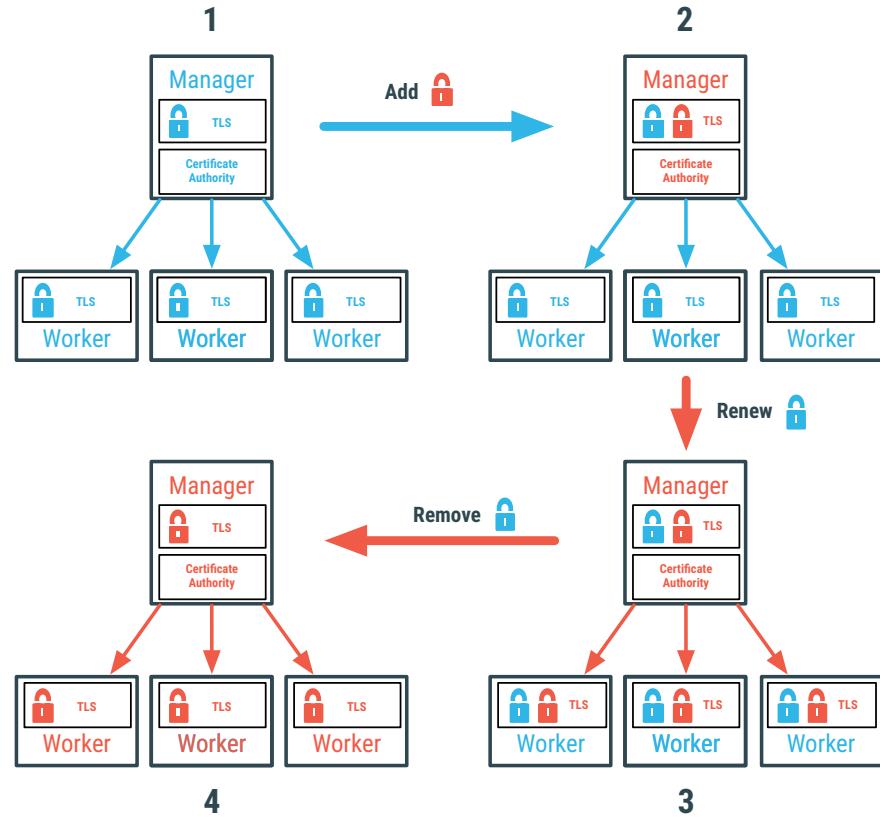
Transparent Root Rotation



3

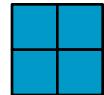
swarmKit

Transparent Root Rotation

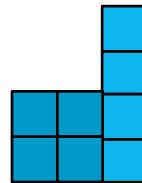


Bringing it all together

Notary for Docker image name resolution



Notary for Docker image name resolution



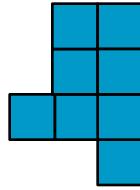
Cryptographically Verified Pulls



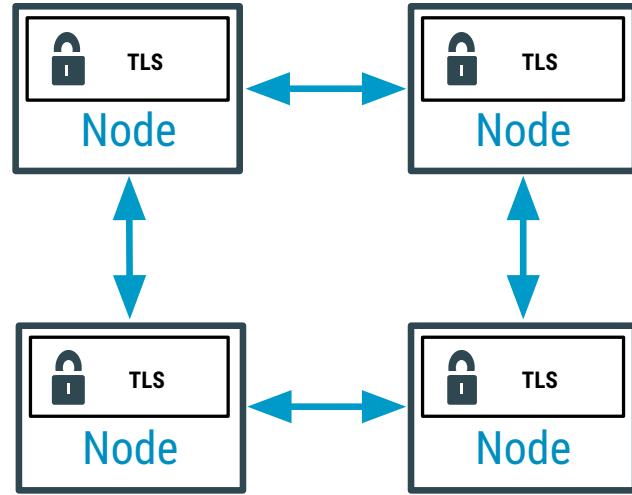
swarmKit delivered Docker containers



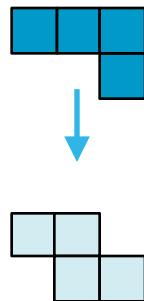
swarmKit delivered Docker containers



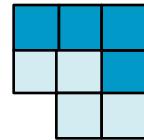
Authorized,
Authenticated,
Encrypted delivery
of Resources



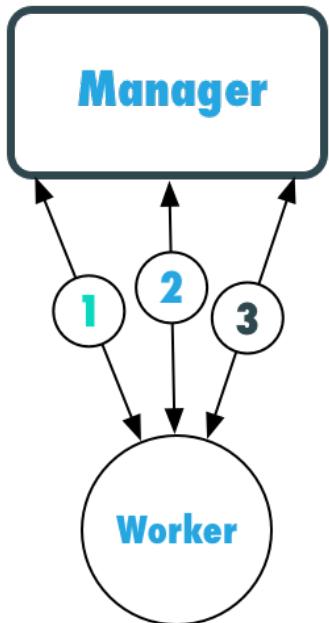
infraKit for swarmKit Bootstrap



infraKit for swarmKit Bootstrap

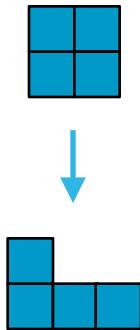


Secure Node Cluster Introduction

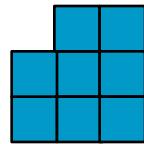


1. Retrieve and validate Root CA Public key material.
2. Submit new CSR along with secret token.
3. Retrieve the signed certificate.

linuxKit as the base OS builder



linuxKit as the base OS builder



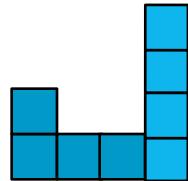
Hardened Configuration



Notary for secure dependency resolution



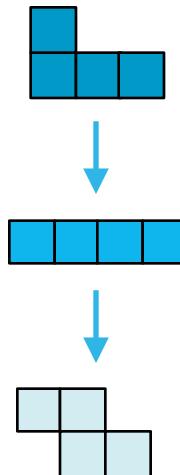
Notary for secure dependency resolution



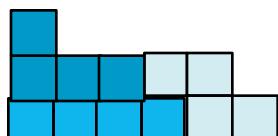
Cryptographically Verified Build



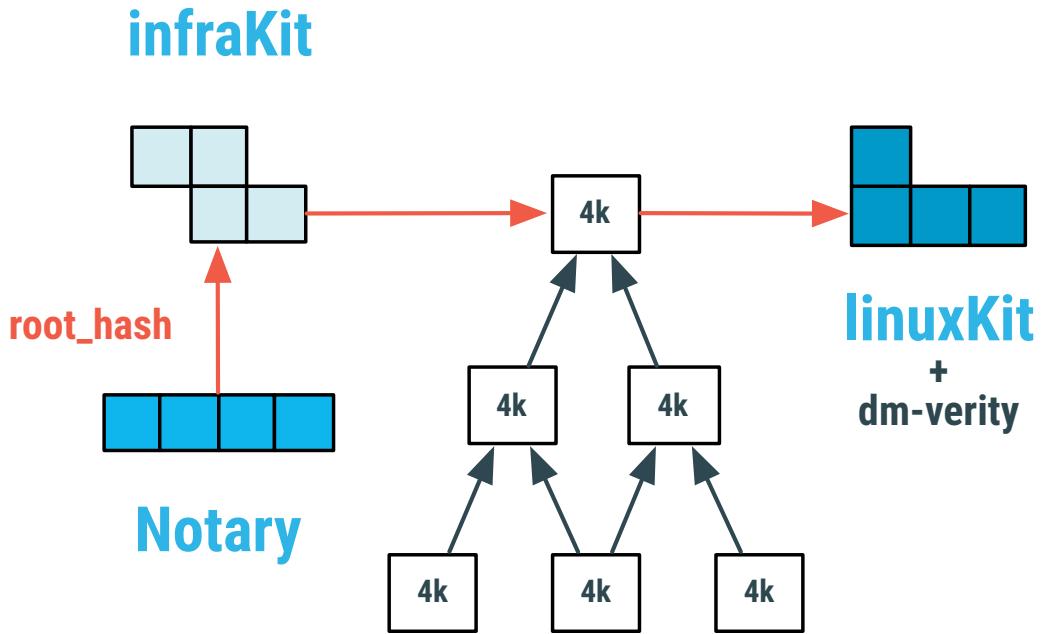
infraKit plus Notary for trusted OS Provisioning

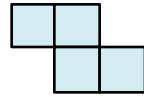


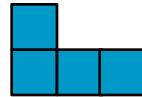
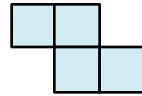
infraKit plus Notary for trusted OS Provisioning

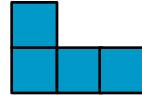
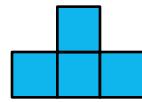
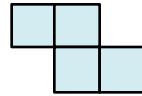


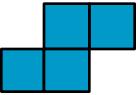
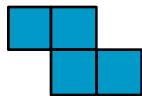
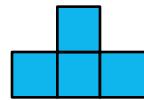
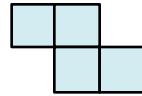
Cryptographically Verified Boot

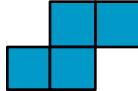
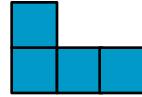
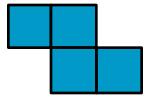
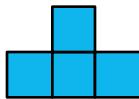
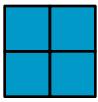
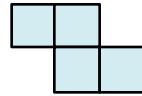


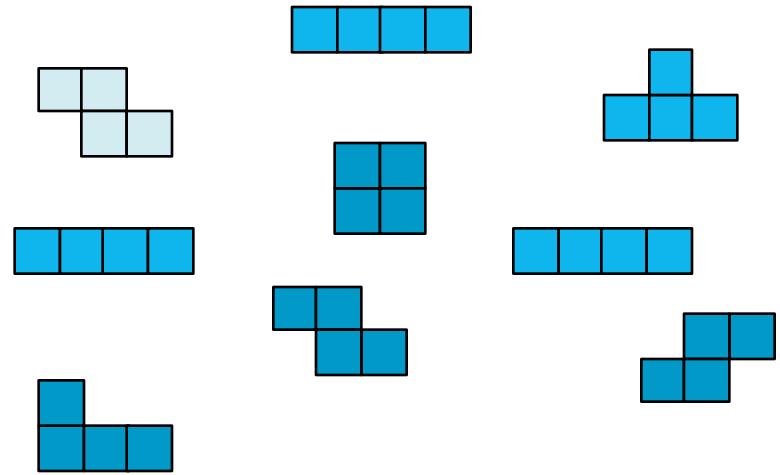


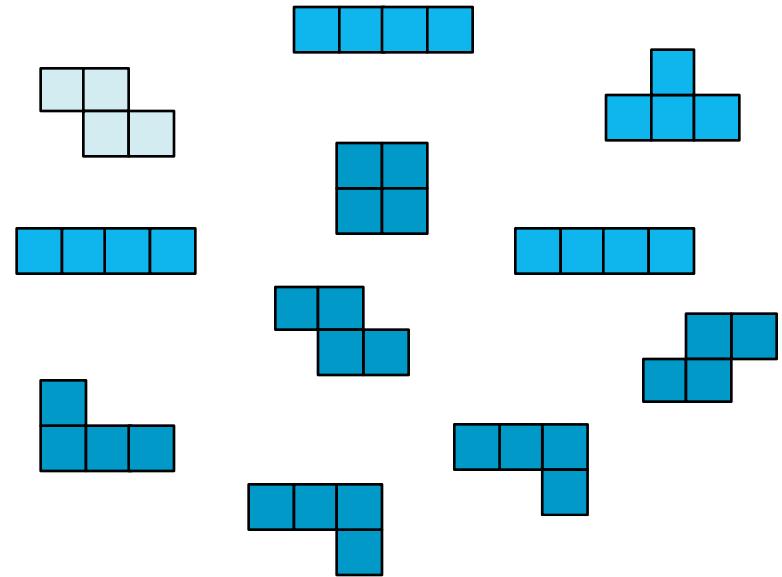


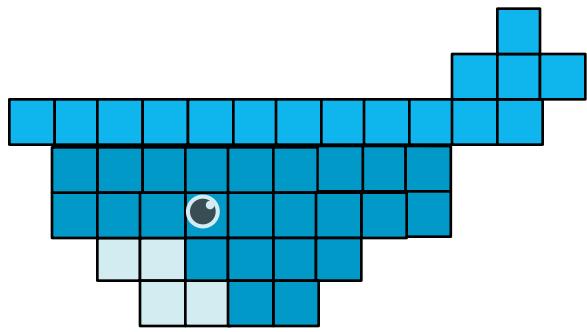


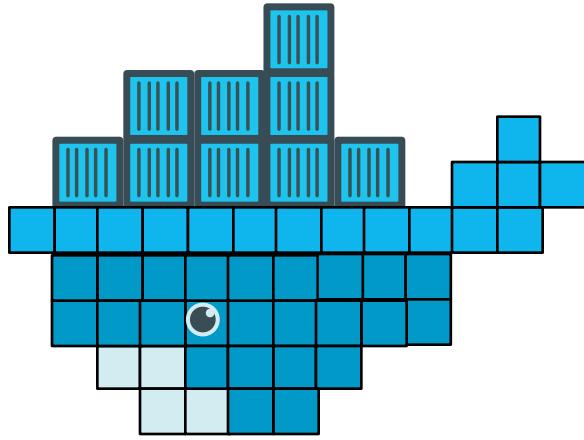












Thank you!